

Article

MisRoBÆRTa: Transformers versus Misinformation

Ciprian-Octavian Truică^{1,2,*}  and Elena-Simona Apostol^{2,*} ¹ InfoLab, Department of Information Technology, Uppsala University, SE-751 05 Uppsala, Sweden² Computer Science and Engineering Department, Faculty of Automatic Control and Computers, University Politehnica of Bucharest, RO-060042 Bucharest, Romania

* Correspondence: ciprian.truica@upb.ro (C.-O.T.); elena.apostol@upb.ro (E.-S.A.)

† These authors contributed equally to this work.

Abstract: Misinformation is considered a threat to our democratic values and principles. The spread of such content on social media polarizes society and undermines public discourse by distorting public perceptions and generating social unrest while lacking the rigor of traditional journalism. Transformers and transfer learning proved to be state-of-the-art methods for multiple well-known natural language processing tasks. In this paper, we propose MisRoBÆRTa, a novel transformer-based deep neural ensemble architecture for misinformation detection. MisRoBÆRTa takes advantage of two state-of-the-art transformers, i.e., BART and RoBERTa, to improve the performance of discriminating between real news and different types of fake news. We also benchmarked and evaluated the performances of multiple transformers on the task of misinformation detection. For training and testing, we used a large real-world news articles dataset (i.e., 100,000 records) labeled with 10 classes, thus addressing two shortcomings in the current research: (1) increasing the size of the dataset from small to large, and (2) moving the focus of fake news detection from binary classification to multi-class classification. For this dataset, we manually verified the content of the news articles to ensure that they were correctly labeled. The experimental results show that the accuracy of transformers on the misinformation detection problem was significantly influenced by the method employed to learn the context, dataset size, and vocabulary dimension. We observe empirically that the best accuracy performance among the classification models that use only one transformer is obtained by BART, while DistilRoBERTa obtains the best accuracy in the least amount of time required for fine-tuning and training. However, the proposed MisRoBÆRTa outperforms the other transformer models in the task of misinformation detection. To arrive at this conclusion, we performed ample ablation and sensitivity testing with MisRoBÆRTa on two datasets.

Keywords: misinformation detection; transformers; benchmark analysis; multi-class classification; large dataset



Citation: Truică, C.-O.; Apostol, E.-S. MisRoBÆRTa: Transformers versus Misinformation. *Mathematics* **2022**, *10*, 569. <https://doi.org/10.3390/math10040569>

Academic Editor: Manuel Vilares-Ferro, Pavel Brazdil and Gaël Dias

Received: 17 January 2022

Accepted: 8 February 2022

Published: 12 February 2022

Publisher's Note: MDPI stays neutral with regard to jurisdictional claims in published maps and institutional affiliations.



Copyright: © 2022 by the authors. Licensee MDPI, Basel, Switzerland. This article is an open access article distributed under the terms and conditions of the Creative Commons Attribution (CC BY) license (<https://creativecommons.org/licenses/by/4.0/>).

1. Introduction

The digital age has seen the emergence of new mass media paradigms for information distribution, substantially different from classical mass media. With new digital-enabled mass media, the communication process is centered around the user, while multimedia content is the new identity of news. Thus, the media landscape has shifted from mass media to personalized social media with these new paradigms. Along with the advantages this progress brings, it aggravates the risk of fake news [1], potentially with detrimental consequences to society, by facilitating the spread of misinformation in the form of propaganda, conspiracy theories, political bias, etc.

Misinformation consists of news articles that are intentionally and verifiably false or inaccurate and which could mislead readers by presenting alleged, imaginary, or false facts about social, economic, and political subjects of interest [2,3]. However, current trends and the accessibility of technology render this proneness even more potentially damaging than it has been historically, having dire consequences to the community and society at

large [4]. Its spread influences communities and creates public polarization regarding elections [5], medical practices [6], etc. Furthermore, the practices of misinforming society are been condemned by politicians and businessmen alike, being considered real threats that undermine democracy and public health [7]; we believe that tools and methods for accurately detecting misinformation are of high relevance to the research community and society at large.

The main motivation for this paper is to evaluate the performance of different transformers for the task of misinformation detection in order to propose a novel deep learning transformer-based model, i.e., MisRoBÆRTa. We used a large real-world news dataset consisting of 100,000 news articles, labeled using 10 different classes, i.e., fake news, satire, extreme bias, hate news, etc. By employing such a dataset, we also address two shortcomings in the current research:

- (1) Increasing the size of the dataset from small to large; and
- (2) Moving the focus of fake news detection from binary to multi-class classification.

The experimental results show that the accuracy of transformers on the multi-class classification task is not influenced by the size of the network, but rather by the method employed to learn the context, and the size of both the dataset and vocabulary. We also include in our benchmark the state-of-the-art FakeBERT model [8].

The research questions for our benchmark are as follows:

- (Q1) Which transformer obtains the overall best accuracy for the task of multi-class classification of misinformation?
- (Q2) Which transformer obtains the overall best runtime performance without the model having a significant decrease in accuracy?
- (Q3) Can there be a balance between accuracy and runtime required for fine-tuning and training?

Our contributions are as follows:

- (1) We propose MisRoBÆRTa, a new transformer-based deep learning architecture for misinformation detection;
- (2) We conducted an in-depth analysis of the current state-of-the-art deep learning, word embeddings, and transformer-based models for the task of fake news and misinformation detection;
- (3) We propose a new balanced dataset for misinformation containing 100,000 articles extracted from the FakeNewsCorpus, where, for each news article, we (1) manually verified its content to make sure that it was correctly labeled; (2) verified the URLs to point to the correct article by matching the titles and authors;
- (4) We performed an in-depth analysis of the proposed dataset;
- (5) We extended the SimpleTransformers package with a multi-class implementation for BART;
- (6) We performed a detailed benchmark analysis using multiple transformers and transfer learning on the task of misinformation and compared the results with the state-of-the-art model FakeBERT [8].

This paper is structured as follows. Section 2 discusses the current research and methods related to misinformation and fake news detection. Section 3 introduces the transformers employed in this study as well as the neural network used for classification. Section 5 presents the datasets and analyzes the experimental results. Section 6 discusses our findings and hints at current challenges that we identify for the task of misinformation detection. Section 7 presents the conclusions and outlines future directions.

2. Related Work

The misinformation and fake news detection problem has been tackled in the literature using different perspectives, from the linguistic approach to more sophisticated natural language processing (NLP) models. Whereas, the linguistic approach attempts to learn language cues that are hard to monitor (e.g., frequencies and patterns of pronouns, con-

junctions, and negative emotion word usage [9]), the NLP approaches build models using content-based features for fake news detection [10].

In article [11], a deep convolutional-based GloVe-enabled model (FNDNet) is used for fake news detection. Although the experimental setup for the tested networks use small epoch numbers, FNDNet shows promising results compared to other deep learning architectures, i.e., classical convolutional neural network (CNN) and long short-term memory (LSTM), both trained using GloVe [12]. Another CNN-based architecture that show promising results is MCNN-TFW (multiple-level convolutional neural network-based) [13]. MCNN-TFW uses semantics features that are determined using a pre-trained Word2Vec [14].

Attention-based models are also good candidates for fake news detection. In [15], a new hierarchical attention network is used for learning GloVe embeddings for different latent aspects of news articles. The solution uses attention weights to support or refute the analyzed claims. The authors emphasize the importance of aspect embeddings for false news prediction. In Trueman et al. [16], the authors present an attention-based convolutional bidirectional LSTM approach to detect fake news. The news articles are classified using six multilevel categories ranging from “true” to “pants on fire”. The authors use the LIAR dataset consisting of 12,836 statements for training, validation, and testing.

Other current approaches use advanced pre-trained model-based detection, i.e., transformers [17–20], or new vectorization techniques, i.e., CLDF [21]. Among the transformers, bidirectional encoder representations from transformers (BERT) (Section 3) has been used the most for fake news detection.

Liu et al. [18] proposed a two-stage BERT-based model for fake news detection. Unlike simple BERT, the authors utilized all hidden states and applied attention mechanisms to calculate weights. The experiments were done using the LIAR dataset containing 12,836 short news statements. Being a relatively small dataset with little information, the proposed model and the compared ones have trouble learning correctly. Furthermore, current research proves that BERT does not provide accurate models if it is fine-tuned with a small dataset or short statements [22].

A comparison in performance between BERT and a deep-learning architecture using CNN and BiLSTM layers is presented in [19]. This article also offers an analysis of the attention coverage of texts for the two models. Although two small (~1000 articles) subsets from the FakeNewsCorpus dataset are used, the classification is not multi-class, as the authors transformed the 11 categories into 2 classes, i.e., real and fake. The results show that, on the two small datasets, BiLSTM-CNN architecture outperforms or has the same accuracy as BERT; thus, strengthening the statement from [22].

In [23], a BERT-RNN architecture uses a relatively big balanced dataset and obtains very good results on a dataset labeled with different levels of veracity i.e., different degrees of truth and falsity. This solution has a hybrid architecture connecting BERT word embeddings with a recurrent neural network (RNN). The RNN network was used to obtain document embeddings. The major issue with this solution is that the results are not explained or discussed, and there is no comparison with other embedding models or deep neural networks.

Another hybrid BERT-based solution is FakeBERT [8]. FakeBERT uses the vectors generated after word-embedding from BERT as input to several 1D-CNN layers. The authors compared their solution with other deep learning solutions, considering a dataset from Kaggle with binary labels. The same dataset is used in FNDNet [11]. By analyzing these two articles written by the same main authors, we can conclude that FakeBERT has slightly better accuracy than FNDNet. Thus, we decided to add FakeBERT to our benchmark.

Other transformers did not get as much attention as BERT, but still, there are some articles that tackle this subject. A weak social supervision RoBERTa-based solution to detect fake news from limited annotated data is proposed in [24]. RoBERTa is an optimized method that improves BERT (Section 3). Their model outperforms other CNN or adversarial neural network-based solutions. In other papers, we found fake news detection solutions based on XLNet, an extension of BERT. In Gautam et al. [25], the authors combined topic information

with an XLNet model. The authors compared their solution (XLNet + LDA) with basic XLNet, and also with BERT on a relatively small annotated dataset (~ 10 K). To achieve an increase in time performance, some solutions use compact language models, such as a lite BERT (ALBERT) [26]. FakeFinder [27] is an ALBERT-based solution for mobile devices that offers real-time fake news detection from Twitter. This small ALBERT-based model achieves comparable F1 scores with a BERT model, although slightly smaller, when tested on a dataset with ~ 2500 tweets and $\sim 117,000$ comments.

To the best of our knowledge, there are almost no benchmarks on misinformation detection using transformers, a multi-class task. The only (found) article that presents a more in-depth benchmark on this subject is [28]. The authors only considered binary classification for five pre-trained models (i.e., BERT, RoBERTa, DistilBERT, ELECTRA, ELMo [29]). The best performance was obtained by RoBERTa, but it is not clear if they used fine-tuning.

After the state-of-the-art analysis, we noticed the following: none of these methods attempted to detect different types of misinformation, e.g., clickbait, extreme bias, but rather focused on the veracity of the articles, binary or multilevel [23], i.e., different degrees of truth and falsity. Moreover, the majority of the solutions used relatively small datasets.

3. Methodology

For our performance evaluation, we used multiple, state-of-the-art transformers that are on multiple natural language processing tasks. For classification, we used a neural network containing the embedding layer and one hidden layer with a dropout layer between them.

3.1. Transformers

For our performance evaluation, we used multiple transformers that employed transfer learning to achieve state-of-the-art performance on multiple natural language processing tasks.

3.1.1. BERT

Bidirectional encoder representations from transformers (BERT) [30], is a deep bidirectional transformer architecture that employs transfer learning to generalize language understanding and can be fine-tuned on the dataset at hand. Classic language models treat textual data as unidirectional or bidirectional sequences of words. Instead, BERT uses the deep bidirectional transformer architecture to learn contextual relations between the words (or sub-words). BERT reads the entire sequence of words at once using the transformer encoder. Thus, the model manages to learn the context of a word based on all of the surrounding words. The experimental results show that the language models built with BERT manage to better determine the language context than the models that treat textual data as sequences of words.

3.1.2. DistilBERT

DistilBERT (Distilled BERT) [31] is a method used to pre-train a smaller general-purpose language representation model by reducing the size of BERT. As BERT, DistilBERT can be fine-tuned with good performances to solve different natural language processing tasks. DistilBERT uses a triple loss combining language modeling, distillation, and cosine-distance losses. The experimental results prove that DistilBERT is a smaller, faster, and lighter model than BERT, which manages to retain the language understanding capabilities.

3.1.3. RoBERTa

A robustly optimized BERT pre-training approach (RoBERTa) [17] is a robustly optimized method that improves BERT's language masking strategy by modifying key hyper-parameters:

- (1) It removes the next-sentence pre-training objective; and
- (2) It increases both the mini-batches and learning rates.

Thus, these modifications improve RoBERTa's masked language modeling objective, leading to better downstream task performance.

3.1.4. DistilRoBERTa

DistilRoBERTa (Distilled RoBERTa) [31,32], similar to DistilBERT, is a method used to pre-train a smaller general-purpose language representation model by reducing the size of RoBERTa. It uses the same methodology employed to distill BERT on the RoBERTa transformer.

3.1.5. ALBERT

A lite BERT (ALBERT) [26] is another optimized method that improves on BERT. It reduces the number of parameters in order to lower memory consumption and increase training speed. The authors of this method show through experimental validation that ALBERT leads to models that scale much better compared to the BERT. Furthermore, ALBERT uses a self-supervised loss that focuses on modeling inter-sentence coherence, which consistently improves downstream tasks with multi-sentence inputs.

3.1.6. DeBERTa

Decoding-enhanced BERT with disentangled attention (DeBERTa) [33] is an improvement of BERT and RoBERTa models. It uses the disentangled attention mechanism, which employs two vectors to represent each word. The first vectors encodes the word's content, while the second it's position. Thus, the attention weights among words are computed using disentangled matrices on their contents and relative positions. Moreover, DeBERTa uses an enhanced mask decoder to replace the output softmax layer to predict the masked tokens for model pre-training.

3.1.7. XLNet

XLNet [34] is an extension of BERT that integrates ideas from the transformer-XL model [35]. It uses an autoregressive method to learn bidirectional contexts by maximizing the expected likelihood over all permutations of the input sequence factorization order. Thus, XLNet overcomes the limitations of BERT through the autoregressive formulation.

3.1.8. ELECTRA

Efficiently learning an encoder that classifies token replacements accurately (ELECTRA) [36] is a pre-training approach used to build two transformer models: the generator and the discriminator. The generator replaces tokens in a sequence and trains as a masked language model. The discriminator predicts which tokens are replaced by the generator in the sequence. We used ELECTRA's discriminator for classification.

3.1.9. XLM

The cross-lingual language model (XLM) [37] is a transformer architecture that employs, during pre-training, either an unsupervised modeling technique that only relies on monolingual data, i.e., casual language modeling (CLM) or masked language modeling (MLM), or a supervised modeling technique that leverages parallel data with a new cross-lingual language model objective, i.e., MLM used in combination with translation language modeling (TLM). CLM models the probability of a word given the previous words in a sentence, while MLM uses the masked language modeling of BERT. TLM is an extension of MLM that concatenates parallel sentences instead of considering monolingual text streams. This model manages to obtain better results than the baseline on multiple tasks, such as cross-lingual classification, unsupervised machine translation, supervised machine translation, low-resource language modeling, and unsupervised cross-lingual word embeddings.

3.1.10. XLM-RoBERTa

XLM-RoBERTa [38] is a transformer-based pre-trained multilingual masked language model. The novelty of this model is that it does not require language tensors to understand which language is used. Thus, the model is able to determine the correct language just from the input IDs.

3.1.11. BART

The bidirectional and autoregressive transformer (BART) [39] uses a standard transformer-based neural machine translation architecture, which can be seen as generalizing BERT. Thus, BART uses a standard sequence-to-sequence for machine translation architecture with a bidirectional encoder, similar to BERT, and a left-to-right decoder, similar to GPT. BART is trained by corrupting text with an arbitrary noising function, and learning a model to reconstruct the original text. During the pre-training task, the model randomly shuffles the order of the original sentences. Moreover, the spans of text are replaced with a single mask token using a novel in-filling scheme. BART is particularly effective when fine-tuned for text generation but it also works well for comprehension tasks [39].

3.2. Classification Network

The network for classification extracts the last hidden layer of the transformer, except for Electra, which uses the discriminator hidden states. All networks use a final dense layer containing perceptrons for classification. Between the transformer model and the classification layer, we used a dropout layer to prevent overfitting. The number of perceptrons in the classification layer is equal to the number of classes. This final hidden layer, added after the transformer model and dropout layer, outputs the label for the multi-class classification. DistilBERT and DistilRoBERTa also use a pre-classification, a ReLU, and a dropout layer between the output of the transformer model and the classification layer. The pre-classification layer size is equal to the dimension of the encoder layers and the pooler layer. The pre-classification layer is used to stabilize the output of the transformer and archive an improved classification performance [40]. For XLM, XLM-RoBERTa, XLNet models, we used a sequence summary layer, which takes a sequence's hidden states and computes a single vector summary, followed by a dense layer for classification.

4. MisRoBÆRTa: Misinformation RoBERTa-BART Ensemble Detection Model

In this section, we propose a new deep learning architecture that develops hierarchical representations of learning using multiple bidirectional long short-term memory (BiLSTM) and convolutional layers (CNN). Figure 1 presents the computational workflow of our novel architecture MisRoBÆRTa: misinformation using RoBERTa-Bart sentence embeddings. In our model, we have two inputs:

- (1) The sentence embeddings constructed with RoBERTa base of size 768; and
- (2) The sentence embeddings extracted with BART, a large of size 1024.

We chose these two transformers for MisRoBÆRTa because they had the best overall results among the pretrained transformers (Section 5) when classification was done using the simple neural network model (Section 3). Each embedding was then the input to a new block (branch) formed using the following layers: BiLSTM, reshape, CNN, max pooling, and flatten. During the ablation testing presented in Section 5.5 we obtained better results when employing two BiLSTM layers for the BART branch. The output of the flatten layers was then concatenated, reshaped, and sent to another block (ensemble branch) containing the same layers that had, as input, the RoBERTa embeddings (RoBERTa branch). Finally, there was a dense layer for the output of the classification.

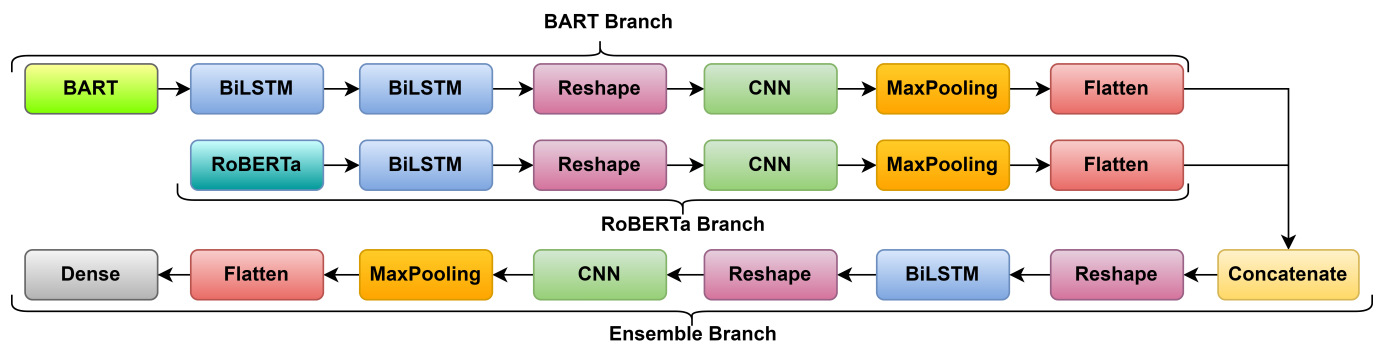


Figure 1. MisRoBÆRTa Architecture.

4.1. MisRoBÆRTa Components

In this subsection, we present MisRoBÆRTa's main deep learning layers: dense, BiLSTM, and CNN. We will not present again RoBERTa and BART, as both transformers are presented in Sections 3.1.3 and 3.1.11, respectively.

4.1.1. Long Short-Term Memory Networks

The long short-term memory network (LSTM) [41] was developed to mitigate the vanishing or the exploding gradient problems that the classical recurrent neural network encountered. By introducing gates within the cell, the LSTM manages to better capture dependencies. The LSTM model uses the *input gate* i_t for deciding what relevant information can be added from the current step t , using the previous hidden state h_t and the input x_t , *forget gate* f_t for deciding which information is used and which is ignored in the current step t , and *output gate* o_t determines the value of the next hidden state h_t . Equation (1) presents of each gate (i.e., f_t , i_t , o_t) the current hidden state (h_{t-1}) and current memory state (c_t), where:

- (1) x_t is the input at time step t ;
- (2) h_t is the output, or next hidden state;
- (3) h_{t-1} is the previous hidden state;
- (4) \tilde{c}_t is the cell input activation vector;
- (5) c_t is the current memory state;
- (6) c_{t-1} is the previous memory state;
- (7) W_f , W_c , W_i , and W_o are the weights for each gate's current input;
- (8) U_f , U_c , U_i , and U_o are the weights for each gate's previous hidden state;
- (9) b_f , b_c , b_i , and b_o are the bias vectors;
- (10) σ_s is the sigmoid activation function;
- (11) σ_h is the hyperbolic tangent activation function;
- (12) \odot operator is the Hadamard product, i.e., the element-wise multiplication function.

$$\begin{aligned}
 i_t &= \sigma_s(W_i x_t + U_i h_{t-1} + b_i) \\
 f_t &= \sigma_s(W_f x_t + U_f h_{t-1} + b_f) \\
 o_t &= \sigma_s(W_o x_t + U_o h_{t-1} + b_o) \\
 \tilde{c}_t &= \sigma_h(W_c x_t + U_c h_{t-1} + b_c) \\
 c_t &= f_t \odot c_{t-1} + i_t * \tilde{c}_t \\
 h_t &= o_t \odot \tanh(c_t)
 \end{aligned} \tag{1}$$

The LSTM processes sequences and it is able to capture *past* information. To also consider *future* information when building the model, the *bidirectional LSTM* (BiLSTM) uses two LSTM with two hidden states $\overrightarrow{h_t}$ and $\overleftarrow{h_t}$ (Equation (2)). The first hidden state $\overrightarrow{h_t}$ processes the input in a forward manner using the past information provided by the

forward LSTM ($\overrightarrow{LSTM_F}$). The second hidden state \overleftarrow{h}_t processes the input in a backwards manner using the future information provided by the backward LSTM ($\overleftarrow{LSTM_B}$).

$$\begin{aligned}\overrightarrow{h}_t &= \overrightarrow{LSTM_F}(x_t) \\ \overleftarrow{h}_t &= \overleftarrow{LSTM_B}(x_t)\end{aligned}\quad (2)$$

Finally, the hidden states \overrightarrow{h}_t and \overleftarrow{h}_t are concatenated at every time-step to enable encoding the information from past and future contexts into one hidden state h'_t (Equation (3)).

$$h'_t = [\overrightarrow{h}_t || \overleftarrow{h}_t] \quad (3)$$

4.1.2. Convolutional Neural Networks

Unlike recurrent neural networks, which are used to detect patterns in sequences, convolutional neural networks (CNNs) were developed to detect patterns in hyperspaces. In text classification, one-dimensional CNNs are used to apply a window of convolution that can extract n-grams. Equation (4) presents a convolution unit r_i , where x_i is a vector of filtered words, W is a weight matrix, b is the bias, and σ is a non-linear activation function.

$$r_i = \sigma(W \cdot x_i + b) \quad (4)$$

4.1.3. Max Pooling

The max pooling mechanism captures the most useful features produced by the previous layer. The mechanism involves sliding a two-dimensional filter over each channel of the feature map, summarizing the features lying within the region covered by the filter. This layer applies an affine transformation to the feature vector to produce logits d . Equation (5) presents the logits d outputted when using a feature map with the dimensions $n_h \times n_w \times n_c$, where n_h , n_w , n_c are the height, the width, and the the number of channels of the feature map, respectively, f is the size of the filter, and s is the stride length.

$$d = \frac{(n_h - f + 1)}{s} \times \frac{(n_w - f + 1)}{s} \times n_c \quad (5)$$

4.1.4. Dense Layer

The dense layer uses perceptron units to classify the output from the previous layer. This layer received the information from the previous layer. Equation (6) presents the representation of the perceptron predicted label \hat{y} , where x_i is the input, w is a weight vector, b represents the bias, and $\delta(\cdot)$ is the activation function. The activation function takes the total input and produces an output for the unit given some threshold. In our model, we use the softmax activation function $\delta(x) = \frac{e^x}{\sum_{i=1}^m e^{x_i}}$ with m the dimension of vector x .

$$\hat{y}_i = \delta(w \cdot x_i + b) \quad (6)$$

A loss function (Equation (7)) is used at each iteration to adjust the weights to align the predicted label \hat{y} to the true label y .

$$L(\hat{y}, y) = -(y \log \hat{y} + (1 - y) \log(1 - \hat{y})) \quad (7)$$

4.2. MisRoBERTa Description

The BiLSTM layers contain 256 units each, 128 units for each long short-term memory (LSTM) and a dropout rate of 0.2. We used the hyperbolic tangent as activation functions and the sigmoid for the recurrent activation function. The CNN layers contain 64 filters, a kernel size of 128, and ReLU as the activation function. After each CNN, we used a max pooling layer to down-sample the output and reduce the number. Flatten was used to transform the two-dimensional output of the max pooling layer into a one-dimensional

tensor. The reshape layers were used to reshape the tensors to be accepted as input for the next layers. Between the BiLSTM and the CNN, we needed a two-dimensional tensor, while between the concatenate layer and the BiLSTM, we needed a three-dimensional tensor. The dense layer was a fully connected layer of perceptrons used for classification. All parameters were determined using hyperparameter tuning.

4.3. MisRoBERTa Architectural Choices

Unidirectional LSTM preserves input information that has passed through its hidden state. Thus, it only uses passed information to make assessments and correct predictions. A BiLSTM instead preserves input information in two ways. As LSTM, the BiLSTM will preserve the information that passes through the hidden layer, i.e., past to future. Unlike LSTM, the BiLSTM also preserves information using the output of the hidden layer, and running backwards the preserved information, i.e., future to past. Thus, BiLSTM combines two hidden states to preserve information from both past and future. BiLSTM is better suited than LSTM to understand context and preserved context. This is also supported by the ablation testing we present in Section 5.5 Furthermore, we observed experimentally that we obtained better results when using two BiLSTM layers instead of one for the BART branch.

After the last BiLSTM layer for each branch, we employed a CNN layer. CNNs have been used successfully for different text classifications and natural language tasks [42]. When using a CNN, the output of each convolution detects patterns of multiple sizes by varying the kernel shapes. These outputs are then concatenated to permit the detection of multiple size patterns of adjacent words. These patterns could be seen as multi-word expressions. Therefore, CNNs can accurately identify and group multi-word expressions in the sentence regardless of their position. These multi-word expressions better represent the context and the representative features of the input text.

5. Experimental Results

5.1. Dataset

The dataset used for experiments was the open-source [FAKENEWSCORPUS](#) (accessed on 6 August 2021) dataset available on GitHub [43]. It is comprised of more than 9.4 million labeled English news articles that contain both textual content (e.g., title, content) and metadata (e.g., publisher, authors). The original dataset is not manually filtered; therefore, some of the labels might not be correct. However, the authors argue that because the corpus is intended to be used in training machine learning algorithms, this shortcoming should not pose a practical issue. Moreover, this issue should help the models better generalize and remove overfitting. For our experiments, we used only the news content and their labels. Moreover, we made sure that the news sources contained articles that belonged to the type given by the class and we manually annotated each news article content and compared it with its original label. The annotation of the textual content was done by computer science students. We used two annotators for each article, introducing a third if there was no consensus between the first two. The annotation task had two main limitations: author bias [44,45] and incomplete annotations [46]. Annotator bias refers to the habits and backgrounds of annotators that label the data and how these can impact the quality of the annotation and, as a result, the model. To mitigate this, we presented the annotators with 10 classes from which to pick only one. Thus the annotator, cannot bring any of his habits within the annotation process. Incomplete annotations refer to the fact that the annotators may not agree among them. In this situation, we added a third annotator who could choose only one of the two labels selected by the previous two annotators. In this way, we mitigated disagreements among the annotators. At the end of this task, we did not find discrepancies between the original labels and the ones determined by the annotators.

Furthermore, we checked that the selected article URLs were pointing to the correct article by matching the titles and authors. We did not apply preprocessing. The final dataset contained 10 classes (currently the 11th class contains 0 articles), with each class represented

by 10,000 documents selected at random (Table 1). We chose this dataset because it is large enough to use with transformers. We also presented the ablation tests for the [Kaggle Fake News dataset](#) (accessed on 9 October 2021) in Section 5.5 used in [8,11].

Table 1. News article classes.

Class	Description
Fake News	Fabricated or distorted information
Satire	Humorous or ironic information
Extreme Bias	Propaganda articles
Conspiracy Theory	Promote conspiracy theories
Junk Science	Scientifically dubious claims
Hate News	Promote discrimination
Clickbait	Credible content, but misleading headline
Proceed With Caution	May be reliable or not
Political	Promote political orientations
Credible	Reliable information

Table 2 presents the dataset's description and token statistics (#Tokens). We observe that the distribution of tokens among the different classes is almost equal. The satire class has the shortest textual content on average, while Hate News has the largest mean number of tokens. We also note that the entire dataset contains 1,337,708 unique tokens and a total of 59,726,182 tokens. By analyzing the dataset, we can affirm that there is no real bias added to the classification problem by the length of the documents.

Table 2. Dataset description and statistics.

Class	Number of Tokens per Document				Number of Tokens per Class	
	Mean	Min	Max	StdDev	Unique	All
Fake News	660.86	7	18,179	899.56	318,808	6,808,099
Satire	245.19	11	4942	232.01	153,509	2,502,701
Extreme Bias	539.71	5	17,402	1192.27	308,827	5,549,908
Conspiracy Theory	800.17	7	17,448	947.65	315,146	8,182,202
Junk Science	511.05	7	12,716	732.69	253,004	5,208,868
Hate News	930.92	5	17,871	2064.67	348,295	9,516,495
Clickbait	361.20	7	5544	361.62	193,844	3,702,766
Unreliable Sources	505.63	7	17,232	982.76	236,495	5,149,370
Political Bias	547.17	9	15,221	795.81	258,269	5,581,168
Credible	737.98	6	15,049	818.02	257,475	7,524,605
Entire dataset statistics	583.99	5	18,179	1037.13	1,337,708	59,726,182

To better understand the token distribution among the different classes, we analyzed documents using unigrams. Table 3 presents the top-10 unigrams extracted using the [NLTK](#) (accessed on 10 October 2021) python package [47]. We extracted the top-10 unigrams for the entire corpus as well as the top-10 unigrams for each class. We used the cosine similarity, based on BERT and FastText, to measure the similarity between the unigrams extracted for the entire corpus and each class. In our implementation, we used [POLYFUZZ](#) (accessed on 14 October 2021) [48]. We observed that the smallest similarity was obtained when comparing the unigrams for the entire corpus with the ones extracted for the satire class, i.e., 0.63 for BERT similarity and 0.44 for FastText similarity. The rest of results are above 0.76 for BERT similarity and 0.67 for FastText similarity. These results show that the terms appearing in each class are well represented within the entire corpus, although for the satire class, are underrepresented. Based on the unigram analysis, we can conclude that there is no real bias added to classes by specific terms, as the distribution of the most frequent terms was similar.

To better understand the context, we used a non-negative matrix factorization (NMF) [49] to extract the top-one topic for the entire dataset and each class. To extract the topics,

we cleaned the corpus removing stopwords and punctuation, then we vectorized the documents using the TFIDF. For our implementation, we used the TFIDF and NMF models from the [SCIKIT-LEARN](#) (accessed on 6 October 2021) [50] library. This analysis was performed in order to determine if the underline hidden context present for each class would influence the behavior of the transformers during fine-tuning. We observed that the best similarity was obtained by the topic extracted for the satire class (Table 4). The other classes had BERT-based similarities over 0.63. These results show that the majority of articles contain context that is present among all classes, which reduces the bias added by the frequent topic and hidden context, while preserving some features that can differentiate between the articles.

Table 3. Top-10 unigram similarity between the entire dataset and each class.

Class	Top-10 Unigrams	Bert Similarity	FastText Similarity
Entire Dataset	People Time Government American World System Year America State Public		
Fake News	people time government world year story market American day God	0.89	0.84
Satire	order close continue user policy agree send deny click advertising	0.63	0.44
Extreme Bias	talk people American government time Russia America state war country	0.87	0.84
Conspiracy	people government time American America world report year country state	0.94	0.91
Junk Science	health people food free time world found body cancer study	0.76	0.67
Hate News	people black white snip percent American whites time country race	0.79	0.71
Clickbait	people Donald time state Clinton government America told country American	0.87	0.84
Unreliable Sources	system Tor operating computer submission stick people public time order	0.79	0.71
Political Bias	people government time American percent year state president tax political	0.89	0.83
Credible	people God Christian government American time world war told Iraq	0.85	0.81

Table 4. Top-1 Topic similarity between the entire dataset and each class.

Class	Top-1 Topic	Bert Similarity	FastText Similarity
Entire Dataset	Banner Preference Navigation Consent Technical Advertising Profile Element User Click		
Fake News	not people Trump year day government time state world no	0.63	0.41
Satire	navigation banner advertising consent technical preference element click profile access	0.97	0.95
Extreme Bias	talk hide link page category template user supply file previous	0.71	0.51
Conspiracy	not report government people president year world state no time	0.65	0.41
Junk Science	food free health reference offer documentary herb nutrition list program	0.68	0.42
Hate News	not black white snip people year percent school no student	0.64	0.41
Clickbait	Trump not president people twitter video state woman year no	0.64	0.41
Unreliable Sources	Tor system tail browser submission computer stick communication GNU bundle	0.70	0.45
Political Bias	Trump not president year state people house government no white	0.64	0.41
Credible	not church people Trump God president no war Bush state	0.63	0.39

Based on the analysis we performed using token statistics, unigram, and topic modeling, we can conclude the following for the dataset. Firstly, the length of the documents and the number of tokens for each class did not add any bias for the models to discriminate based on these two dimensions. Secondly, the unique tokens per class versus the unique tokens for the entire dataset shows that some tokens were representative for some classes. Thirdly, the most frequent unigrams did not influence the models, as the similarity between the unigrams extracted for each class and for the entire corpus was high. Finally, the global context for each class had a moderate similarity with the context for the entire corpus. These results show that, among the different classes, there were some clear discriminative features.

5.2. Fine-Tuning

We used the [SIMPLETRANSFORMERS](#) (accessed on 12 October 2021) [51] library and the pre-trained transformers from [HUGGINGFACE](#) (accessed on 9 October 2021) [52]. We extended SIMPLETRANSFORMERS with our own implementation for a BART-based multi-class classifier. Table 5 presents the transformers employed for the experiments.

The ALBERT v2 models did not have dropout. We used the ELECTRA’s discriminator for classification.

Table 5. Transformer details.

Model	Pretrained Model Name	Encoder Layers	Hidden State	Attention Heads	Parameters
BERT base	bert-base-cased	12	768	12	110 M
BERT large	bert-large-cased	24	1024	16	335 M
DistilBERT	distilbert-base-cased	6	768	12	65 M
RoBERTa base	roberta-base	12	768	12	125 M
RoBERTa large	roberta-large	24	1024	16	355 M
DistilRoBERTa	distilroberta-base	6	768	12	82 M
XLNet base	xlnet-base-cased	12	768	12	110 M
XLNet large	xlnet-large-cased	24	1024	16	340 M
ALBERT base v1	albert-base-v1	12	768	12	11 M
ALBERT base v2	albert-base-v2	12	768	12	11 M
ALBERT xxlarge v1	albert-xxlarge-v1	12	4096	64	223 M
ALBERT xxlarge v2	albert-xxlarge-v2	12	4096	64	223 M
DeBERTa base	microsoft/deberta-base	12	768	12	140 M
DeBERTa large	microsoft/deberta-large	24	1024	16	400 M
ELECTRA base	google/electra-base-discriminator	12	768	12	110 M
ELECTRA large	google/electra-large-discriminator	24	1024	16	335 M
XLM	xlm-mlm-100-1280	16	1280	16	550 M
XLM-RoBERTa base	xlm-roberta-base	12	768	8	270 M
XLM-RoBERTa large	lm-roberta-large	24	1024	16	550 M
BART base	facebook/bart-base	12	768	16	139 M
BART large	facebook/bart-large	12	1024	16	406 M

The employed transformers were fine-tuned by training them on our dataset for 10 epochs with an early stopping mechanism that evaluated the loss over 5 epochs. We split the dataset using a 64%–16%–20% training–validation–testing ratio with random seeding. For each split, we maintained the stratification of the labels. We applied this method for 10 rounds of tests.

5.3. Classification Results

We used an NVIDIA® DGX Station™ containing 4 NVIDIA® V100 Tensor Core GPUs for our experiments. Table 6 presents the classification results. We observed that all transformer models had an accuracy of over 85%, with the lowest accuracy obtained by ALBERT base v2. For comparison reasons, we also tested the state-of-the-art model FakeBERT [8]. As the authors of this model did not provide their implementation, we implemented FakeBERT using BERT base together with the [TENSORFLOW](#) (accessed on 8 October 2021) [53] and [KERAS](#) (accessed on 8 October 2021) [54] libraries. The code for the MisRoBERTa, transformer implementation, and FakeBERT are publicly available on GitHub at https://github.com/cipriantruica/MisRoBERTa_Transformers-vs-misinformation (accessed on 8 October 2021).

The BERT, RoBERTa, DeBERTa, XLNet, ELECTRA, and XLM-RoBERTa base models obtained better results than their counterpart large models. For these models, we can conclude that a smaller number of encoder layers with smaller sizes better generalized and managed to correctly extract the hidden context within the texts.

The reverse happened for ALBERT and BART, where the large models outperformed the base models. Although the difference in performance between BART base and BART large was small ($\sim 0.60\%$) for ALBERT, the difference was between $\sim 1.30\%$ and $\sim 2.3\%$. For BART, we can conclude that the autoregressive method, the sequence-to-sequence architecture, and the left–right decoder played an important factor and improved performance as the model’s number of layers and size increased. For ALBERT, the self-supervised loss that focused on modeling inter-sentence coherence improved the model’s performance as the

network size increased. Furthermore, the use of dropout improved ALBERT performance significantly, as ALBERT v2 versions had the overall worst performance.

We observe that DeBERTa obtained slightly better results than BERT, regardless if it was the base or the large model. DeBERTa base obtained an accuracy slightly worse than the RoBERTa base, while DeBERTa large outperformed RoBERTa large. XLM obtained the second worst accuracy among the tested transformers.

Table 6. Misinformation classification results (note: italic text marks similar results while bold text shows the overall best result).

Model	Accuracy	Micro Precision	Macro Precision	Micro Recall	Macro Recall	Execution Time (Hours)
MisRoBERTa	92.50 ± 0.26	92.50 ± 0.26	92.69 ± 0.21	92.50 ± 0.26	92.50 ± 0.26	1.83 ± 0.01
BERT base	90.03 ± 0.19	90.03 ± 0.19	90.05 ± 0.21	90.03 ± 0.19	90.03 ± 0.19	5.12 ± 0.22
BERT large	89.12 ± 0.14	89.12 ± 0.14	89.11 ± 0.15	89.12 ± 0.14	89.12 ± 0.14	8.97 ± 0.05
DistilBERT	89.36 ± 0.15	89.36 ± 0.15	89.40 ± 0.16	89.36 ± 0.15	89.36 ± 0.15	2.30 ± 0.01
RoBERTa base	<i>91.36 ± 0.15</i>	91.36 ± 0.15	91.39 ± 0.16	91.36 ± 0.15	91.36 ± 0.15	3.95 ± 0.15
RoBERTa large	88.60 ± 0.23	88.61 ± 0.22	88.85 ± 0.23	88.59 ± 0.24	88.62 ± 0.21	6.55 ± 0.32
DistilRoBERTa	<i>91.32 ± 0.10</i>	91.32 ± 0.10	91.34 ± 0.08	91.32 ± 0.10	91.32 ± 0.10	1.93 ± 0.01
ALBERT base v1	86.92 ± 0.17	86.92 ± 0.17	86.95 ± 0.18	86.92 ± 0.17	86.92 ± 0.17	2.16 ± 0.19
ALBERT base v2	85.05 ± 0.20	85.05 ± 0.20	85.08 ± 0.17	85.05 ± 0.20	85.05 ± 0.20	2.76 ± 0.17
ALBERT xxlarge v1	89.20 ± 0.04	89.20 ± 0.04	89.36 ± 0.02	89.20 ± 0.04	89.20 ± 0.04	9.98 ± 0.09
ALBERT xxlarge v2	86.51 ± 2.23	86.51 ± 2.23	86.66 ± 2.21	86.51 ± 2.23	86.51 ± 2.23	11.14 ± 0.04
DeBERTa base	90.56 ± 0.08	90.56 ± 0.08	90.56 ± 0.10	90.56 ± 0.08	90.56 ± 0.08	6.90 ± 0.35
DeBERTa large	89.93 ± 0.27	89.93 ± 0.27	89.96 ± 0.29	89.93 ± 0.27	89.93 ± 0.27	11.06 ± 1.10
XLNet base	89.94 ± 0.09	89.94 ± 0.09	89.94 ± 0.09	89.94 ± 0.09	89.94 ± 0.09	5.50 ± 0.06
XLNet large	88.05 ± 0.39	88.04 ± 0.38	88.38 ± 0.39	88.03 ± 0.37	88.07 ± 0.40	10.44 ± 1.54
ELECTRA base	87.10 ± 0.24	87.10 ± 0.24	87.12 ± 0.22	87.10 ± 0.24	87.10 ± 0.24	4.49 ± 0.34
ELECTRA large	86.92 ± 0.09	86.92 ± 0.09	86.92 ± 0.08	86.92 ± 0.09	86.92 ± 0.09	8.74 ± 0.43
XLM	85.82 ± 0.19	85.82 ± 0.19	85.90 ± 0.23	85.82 ± 0.19	85.82 ± 0.19	5.16 ± 0.01
XLM-RoBERTa base	89.78 ± 0.17	89.78 ± 0.17	89.77 ± 0.17	89.78 ± 0.17	89.78 ± 0.17	6.97 ± 0.38
XLM-RoBERTa large	87.50 ± 0.61	87.50 ± 0.61	87.58 ± 0.58	87.50 ± 0.61	87.50 ± 0.61	9.46 ± 2.16
BART base	<i>91.35 ± 0.04</i>	91.35 ± 0.04	91.35 ± 0.04	91.35 ± 0.04	91.35 ± 0.04	4.17 ± 0.16
BART large	<i>91.94 ± 0.15</i>	91.94 ± 0.15	91.97 ± 0.16	91.94 ± 0.15	91.94 ± 0.15	6.79 ± 0.29
FakeBERT [8]	70.18 ± 0.01	70.18 ± 0.01	70.21 ± 0.03	70.18 ± 0.01	70.18 ± 0.01	2.59 ± 0.01

Both DistilBERT and DistilRoBERTa were outperformed by BERT, respectively, RoBERTa base models. The difference in accuracy between the base and distilled models was very small: $\sim 0.04\%$ between DistilBERT and BERT base and $\sim 0.70\%$ between DistilRoBERTa and RoBERTa base. Although DistilBERT outperformed BERT large with a small increase in accuracy ($\sim 0.24\%$), DistilRoBERTa outperformed RoBERTa large with a large increase in accuracy ($\sim 2.72\%$). We can conclude that, for the use case of misinformation classification, distillation is not a good approach if the computational and memory resources needed for the base approaches are available.

The micro and macro recall metrics are the most relevant for misinformation classification, as they show the correctly classified news articles relative to all the news articles, regardless of the predicted label. The micro and macro precisions manage to show the percentage of the articles, which are relevant. All transformers obtained results over 85% for these four metrics, meaning that they managed to generalize well for the task of misinformation detection.

The overall best performing model was the novel MisRoBERTa with an accuracy of 92.50%, while the overall best performing transformer model was BART large, with an accuracy of 91.94%. The next three transformers that obtained accuracies over 91% were RoBERTa base (91.36%), BART base (91.35%), and DistilRoBERTa 91.32%.

We observed that the state-of-the-art-model FakeBERT had very low performance on our dataset. Although the original paper shows that FakeBERT obtained an accuracy of 98.9% on a binary dataset from Kaggle, in our multi-class dataset, it obtained only $\sim 70\%$. These results might be a consequence of the following:

- (1) FakeBERT is developed for binary classification, while we used it for multi-class classification;
- (2) To change FakeBERT from a binary classifier to a multi-class classifier, we modified the number of units of its final DENSE layer from 2 to the number of classes, i.e., 10;
- (3) We only trained FakeBERT for 10 epochs, to be consistent in our performance experiments (note: this is also the number of epochs used in the original paper).

MisRoBÆRTa had the overall best classification results with an accuracy of 92.50%. This accuracy was a direct result of the way context was preserved through the use of:

- (1) BART large and RoBERTa base;
- (2) BiLSTM layers that combined two hidden states to preserve information from both past and future;
- (3) CNNs to generate multi-word expressions, which better represent the feature space.

5.4. Runtime Evaluation

To evaluate the runtime, we ran each transformer for 10 epochs without early stopping. As already mentioned, we used an NVIDIA® DGX Station™ containing 4 NVIDIA® V100 Tensor Core GPUs for our evaluation. Figure 2 presents the mean execution time and the standard deviation for 10 rounds of tests. As expected, the large models are much slower than their base counterparts. Moreover, we note that the distilled versions DistilRoBERTa, ALBERT base, and DistilBERT outperformed both BERT and RoBERTa. MisRoBÆRTa had the fastest runtime among the tested models managing to converge in 1.83 h (this included the maximum time taken to create the sentence embeddings in parallel using the two transformers: 1.53 h for BART and 0.35 h for RoBERTa).

For the base models, the fastest transformer model was DistilRoBERTa, which required, on average, 1.93 h for fine-tuning and training, while the slowest was XLM-RoBERTa, which finished training and fine-tuning in a little less than 7 h on average. For the large models, RoBERTa large was the fastest, requiring on average 6.55 h for fine-tuning and training, while the slowest was ALBERT large v2, which finished in 11.14 h on average. FakeBERT converged in 2.59 h on average. This result was directly impacted by the CNN and max pooling layers it employed to minimize the input space. XLM runtime performance at 5.16 h on average was between BERT base and XLNet base. In conclusion, the fastest transformer model that also achieved a high accuracy was DistilRoBERTa.

5.5. Ablation Testing

In Tables 7 and 8, we present ablation and sensitivity testing for MisRoBÆRTa using the publicly available [Kaggle Fake News Dataset](#) and the FakeNewsCorpus dataset presented in Section 5. The tests include (1) testing with LSTM and BiLSTM units, (2) variations on the number of layers for each branch, e.g., 1, 2, and 3 layers, and (3) variation in the number of units for each LSTM layer, e.g., 32, 64, 128. The results are aggregated after 10 runs, for each execution, the dataset was split at random without replacement, but keeping the class ratio into 64% training, 16% validation, and 20% testing. We present only the network's training time, the sentence embedding creation time was not taken into account. We should note that the results we obtained during ablation and sensitivity testing on the publicly available [Kaggle Fake News dataset](#) are similar to the current state-of-the-art results obtained on this corpus by [8,11].

Table 7. MisRoBÆRTa ablation and sensitivity testing using the [Kaggle Fake News Dataset](#). (Note: bold text marks the overall best result).

32 Units/Layer			MisRoBÆRTa with LSTM Cells						MisRoBÆRTa with BiLSTM Cells					
BART Branch	RoBERTa Branch	Ensemble Branch	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)
1 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	97.10 ± 0.43	97.10 ± 0.43	97.16 ± 0.42	97.10 ± 0.43	97.07 ± 0.43	0.01 ± 0.00	97.35 ± 0.37	97.35 ± 0.37	97.40 ± 0.36	97.35 ± 0.37	97.33 ± 0.37	0.01 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	97.14 ± 0.35	97.14 ± 0.35	97.17 ± 0.35	97.14 ± 0.35	97.12 ± 0.35	0.02 ± 0.00	97.30 ± 0.42	97.30 ± 0.42	97.33 ± 0.38	97.30 ± 0.42	97.29 ± 0.45	0.03 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	97.10 ± 0.57	97.10 ± 0.57	97.13 ± 0.56	97.10 ± 0.57	97.08 ± 0.57	0.03 ± 0.00	97.32 ± 0.37	97.32 ± 0.37	97.34 ± 0.35	97.32 ± 0.37	97.34 ± 0.36	0.05 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	97.06 ± 0.28	97.06 ± 0.28	97.09 ± 0.28	97.06 ± 0.28	97.05 ± 0.27	0.04 ± 0.00	97.02 ± 0.48	97.02 ± 0.48	97.10 ± 0.44	97.02 ± 0.48	96.99 ± 0.49	0.06 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	97.15 ± 0.35	97.15 ± 0.35	97.16 ± 0.35	97.15 ± 0.35	97.15 ± 0.35	0.05 ± 0.00	97.43 ± 0.28	97.43 ± 0.28	97.44 ± 0.26	97.43 ± 0.28	97.43 ± 0.30	0.08 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	97.16 ± 0.37	97.16 ± 0.37	97.18 ± 0.37	97.16 ± 0.37	97.17 ± 0.35	0.07 ± 0.00	97.36 ± 0.32	97.36 ± 0.32	97.38 ± 0.30	97.36 ± 0.32	97.38 ± 0.33	0.11 ± 0.00
1 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	97.09 ± 0.46	97.09 ± 0.46	97.10 ± 0.45	97.09 ± 0.46	97.11 ± 0.44	0.08 ± 0.00	97.33 ± 0.32	97.33 ± 0.32	97.35 ± 0.30	97.33 ± 0.32	97.32 ± 0.34	0.13 ± 0.00
1 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	97.07 ± 0.37	97.07 ± 0.37	97.08 ± 0.38	97.07 ± 0.37	97.08 ± 0.36	0.09 ± 0.00	97.14 ± 0.41	97.14 ± 0.41	97.16 ± 0.39	97.14 ± 0.41	97.17 ± 0.39	0.15 ± 0.00
1 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	97.30 ± 0.37	97.30 ± 0.37	97.35 ± 0.38	97.30 ± 0.37	97.27 ± 0.36	0.11 ± 0.00	97.24 ± 0.53	97.24 ± 0.53	97.31 ± 0.48	97.24 ± 0.53	97.22 ± 0.55	0.18 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	97.56 ± 0.29	97.56 ± 0.29	97.57 ± 0.27	97.56 ± 0.29	97.55 ± 0.30	0.12 ± 0.00	97.57 ± 0.29	97.57 ± 0.29	97.58 ± 0.28	97.57 ± 0.29	97.57 ± 0.31	0.19 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	97.34 ± 0.36	97.34 ± 0.36	97.38 ± 0.35	97.34 ± 0.36	97.32 ± 0.36	0.13 ± 0.00	97.44 ± 0.38	97.44 ± 0.38	97.44 ± 0.38	97.44 ± 0.38	97.44 ± 0.37	0.21 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	97.20 ± 0.38	97.20 ± 0.38	97.24 ± 0.36	97.20 ± 0.38	97.18 ± 0.39	0.15 ± 0.00	97.31 ± 0.44	97.31 ± 0.44	97.34 ± 0.42	97.31 ± 0.44	97.30 ± 0.46	0.23 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	97.31 ± 0.47	97.31 ± 0.47	97.38 ± 0.41	97.31 ± 0.47	97.28 ± 0.50	0.16 ± 0.00	97.41 ± 0.37	97.41 ± 0.37	97.42 ± 0.38	97.41 ± 0.37	97.41 ± 0.37	0.25 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	97.44 ± 0.26	97.44 ± 0.26	97.45 ± 0.26	97.44 ± 0.26	97.42 ± 0.27	0.18 ± 0.00	97.26 ± 0.37	97.26 ± 0.37	97.30 ± 0.36	97.26 ± 0.37	97.25 ± 0.36	0.28 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	97.30 ± 0.33	97.30 ± 0.33	97.34 ± 0.25	97.30 ± 0.33	97.29 ± 0.35	0.19 ± 0.00	97.40 ± 0.44	97.40 ± 0.44	97.44 ± 0.42	97.40 ± 0.44	97.37 ± 0.44	0.30 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	97.37 ± 0.41	97.37 ± 0.41	97.40 ± 0.41	97.37 ± 0.41	97.36 ± 0.41	0.21 ± 0.00	97.40 ± 0.37	97.40 ± 0.37	97.42 ± 0.36	97.40 ± 0.37	97.39 ± 0.37	0.33 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	97.39 ± 0.36	97.39 ± 0.36	97.39 ± 0.37	97.39 ± 0.36	97.39 ± 0.36	0.22 ± 0.00	97.46 ± 0.22	97.46 ± 0.22	97.47 ± 0.20	97.46 ± 0.22	97.47 ± 0.20	0.35 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	97.09 ± 0.35	97.09 ± 0.35	97.14 ± 0.34	97.09 ± 0.35	97.10 ± 0.33	0.24 ± 0.00	97.16 ± 0.37	97.16 ± 0.37	97.17 ± 0.38	97.16 ± 0.37	97.16 ± 0.36	0.38 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	97.43 ± 0.39	97.43 ± 0.39	97.46 ± 0.37	97.43 ± 0.39	97.41 ± 0.41	0.26 ± 0.00	97.45 ± 0.39	97.45 ± 0.39	97.48 ± 0.39	97.45 ± 0.39	97.44 ± 0.39	0.40 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	97.42 ± 0.33	97.42 ± 0.33	97.44 ± 0.34	97.42 ± 0.33	97.40 ± 0.32	0.27 ± 0.01	97.36 ± 0.35	97.36 ± 0.35	97.40 ± 0.35	97.36 ± 0.35	97.35 ± 0.35	0.43 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	97.19 ± 0.29	97.19 ± 0.29	97.30 ± 0.23	97.19 ± 0.29	97.15 ± 0.32	0.29 ± 0.00	97.41 ± 0.27	97.41 ± 0.27	97.42 ± 0.29	97.41 ± 0.27	97.40 ± 0.27	0.46 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	97.25 ± 0.44	97.25 ± 0.44	97.28 ± 0.41	97.25 ± 0.44	97.25 ± 0.42	0.31 ± 0.01	97.33 ± 0.47	97.33 ± 0.47	97.38 ± 0.44	97.33 ± 0.47	97.33 ± 0.43	0.48 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	97.31 ± 0.32	97.31 ± 0.32	97.36 ± 0.28	97.31 ± 0.32	97.29 ± 0.34	0.32 ± 0.01	97.30 ± 0.39	97.30 ± 0.39	97.34 ± 0.32	97.30 ± 0.39	97.29 ± 0.42	0.51 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	97.23 ± 0.41	97.23 ± 0.41	97.26 ± 0.40	97.23 ± 0.41	97.23 ± 0.38	0.34 ± 0.01	97.48 ± 0.42	97.48 ± 0.42	97.48 ± 0.42	97.48 ± 0.42	97.50 ± 0.41	0.54 ± 0.01
3 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	97.41 ± 0.34	97.41 ± 0.34	97.43 ± 0.35	97.41 ± 0.34	97.39 ± 0.32	0.36 ± 0.01	97.31 ± 0.30	97.31 ± 0.30	97.35 ± 0.31	97.31 ± 0.30	97.30 ± 0.30	0.57 ± 0.01
3 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	97.24 ± 0.29	97.24 ± 0.29	97.29 ± 0.31	97.24 ± 0.29	97.21 ± 0.28	0.38 ± 0.01	97.42 ± 0.45	97.42 ± 0.45	97.44 ± 0.43	97.42 ± 0.45	97.42 ± 0.46	0.60 ± 0.02
3 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	97.15 ± 0.31	97.15 ± 0.31	97.17 ± 0.30	97.15 ± 0.31	97.13 ± 0.32	0.40 ± 0.01	97.30 ± 0.44	97.30 ± 0.44	97.32 ± 0.43	97.30 ± 0.44	97.31 ± 0.43	0.63 ± 0.02
64 Units/Layer			MisRoBÆRTa with LSTM Cells						MisRoBÆRTa with BiLSTM Cells					
BART Branch	RoBERTa Branch	Ensemble Branch	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)
1 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	97.24 ± 0.78	97.24 ± 0.78	97.29 ± 0.72	97.24 ± 0.78	97.25 ± 0.77	0.01 ± 0.00	97.42 ± 0.47	97.42 ± 0.47	97.49 ± 0.40	97.42 ± 0.47	97.39 ± 0.48	0.01 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	97.36 ± 0.42	97.36 ± 0.42	97.41 ± 0.39	97.36 ± 0.42	97.34 ± 0.43	0.02 ± 0.00	97.29 ± 0.48	97.29 ± 0.48	97.35 ± 0.40	97.29 ± 0.48	97.29 ± 0.50	0.03 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	97.42 ± 0.47	97.42 ± 0.47	97.44 ± 0.41	97.42 ± 0.47	97.41 ± 0.48	0.03 ± 0.00	97.20 ± 0.50	97.20 ± 0.50	97.25 ± 0.46	97.20 ± 0.50	97.20 ± 0.48	0.04 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	97.48 ± 0.37	97.48 ± 0.37	97.48 ± 0.37	97.48 ± 0.37	97.48 ± 0.38	0.04 ± 0.00	97.26 ± 0.28	97.26 ± 0.28	97.36 ± 0.17	97.26 ± 0.28	97.24 ± 0.33	0.06 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	97.33 ± 0.39	97.33 ± 0.39	97.34 ± 0.39	97.33 ± 0.39	97.33 ± 0.38	0.05 ± 0.00	97.13 ± 0.39	97.13 ± 0.39	97.25 ± 0.30	97.13 ± 0.39	97.10 ± 0.43	0.08 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	97.36 ± 0.42	97.36 ± 0.42	97.38 ± 0.40	97.36 ± 0.42	97.35 ± 0.40	0.06 ± 0.00	97.35 ± 0.39	97.35 ± 0.39	97.40 ± 0.37	97.35 ± 0.39	97.34 ± 0.39	0.10 ± 0.00
1 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	97.32 ± 0.46	97.32 ± 0.46	97.35 ± 0.45	97.32 ± 0.46	97.31 ± 0.46	0.07 ± 0.00	97.29 ± 0.77	97.29 ± 0.77	97.37 ± 0.67	97.29 ± 0.77	97.29 ± 0.73	0.12 ± 0.00
1 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	97.53 ± 0.34	97.53 ± 0.34	97.56 ± 0.33	97.53 ± 0.34	97.50 ± 0.35	0.09 ± 0.00	97.20 ± 0.79	97.20 ± 0.79	97.30 ± 0.66	97.20 ± 0.79	97.22 ± 0.73	0.14 ± 0.00
1 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	97.32 ± 0.40	97.32 ± 0.40	97.32 ± 0.39	97.32 ± 0.40	97.33 ± 0.38	0.10 ± 0.00	97.33 ± 0.50	97.33 ± 0.50	97.38 ± 0.45	97.33 ± 0.50	97.35 ± 0.47	0.16 ± 0.00
2 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	97.84 ± 0.50	97.84 ± 0.50	97.85 ± 0.50	97.84 ± 0.50	97.84 ± 0.48	0.11 ± 0.00	97.85 ± 0.23	97.85 ± 0.23	97.88 ± 0.24	97.85 ± 0.23	97.83 ± 0.22	0.18 ± 0.00

Table 7. Cont.

64 Units/Layer			MisRoBÆRTa with LSTM Cells						MisRoBÆRTa with BiLSTM Cells					
BART Branch	RoBERTa Branch	Ensemble Branch	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)
2 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	97.56 ± 0.38	97.56 ± 0.38	97.58 ± 0.40	97.56 ± 0.38	97.55 ± 0.37	0.12 ± 0.00	97.40 ± 0.38	97.40 ± 0.38	97.51 ± 0.34	97.40 ± 0.38	97.36 ± 0.38	0.20 ± 0.00
2 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	97.56 ± 0.48	97.56 ± 0.48	97.60 ± 0.42	97.56 ± 0.48	97.56 ± 0.50	0.14 ± 0.00	97.39 ± 0.36	97.39 ± 0.36	97.45 ± 0.37	97.39 ± 0.36	97.37 ± 0.33	0.22 ± 0.00
2 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	97.53 ± 0.48	97.53 ± 0.48	97.57 ± 0.48	97.53 ± 0.48	97.50 ± 0.47	0.15 ± 0.00	97.33 ± 0.46	97.33 ± 0.46	97.41 ± 0.35	97.33 ± 0.46	97.32 ± 0.50	0.24 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	97.56 ± 0.45	97.56 ± 0.45	97.61 ± 0.42	97.56 ± 0.45	97.55 ± 0.43	0.16 ± 0.00	97.38 ± 0.65	97.38 ± 0.65	97.42 ± 0.60	97.38 ± 0.65	97.38 ± 0.63	0.26 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	97.53 ± 0.38	97.53 ± 0.38	97.56 ± 0.36	97.53 ± 0.38	97.52 ± 0.39	0.18 ± 0.00	97.35 ± 0.41	97.35 ± 0.41	97.43 ± 0.34	97.35 ± 0.41	97.34 ± 0.41	0.28 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	97.52 ± 0.35	97.52 ± 0.35	97.55 ± 0.34	97.52 ± 0.35	97.52 ± 0.35	0.19 ± 0.00	97.47 ± 0.21	97.47 ± 0.21	97.52 ± 0.20	97.47 ± 0.21	97.44 ± 0.22	0.30 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	97.52 ± 0.37	97.52 ± 0.37	97.57 ± 0.34	97.52 ± 0.37	97.50 ± 0.39	0.21 ± 0.00	97.46 ± 0.34	97.46 ± 0.34	97.52 ± 0.29	97.46 ± 0.34	97.47 ± 0.31	0.33 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	97.63 ± 0.31	97.63 ± 0.31	97.65 ± 0.30	97.63 ± 0.31	97.60 ± 0.31	0.22 ± 0.01	97.46 ± 0.25	97.46 ± 0.25	97.49 ± 0.18	97.46 ± 0.25	97.46 ± 0.28	0.36 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	97.72 ± 0.25	97.72 ± 0.25	97.73 ± 0.26	97.72 ± 0.25	97.70 ± 0.25	0.24 ± 0.01	97.60 ± 0.37	97.60 ± 0.37	97.65 ± 0.40	97.60 ± 0.37	97.56 ± 0.37	0.38 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	97.67 ± 0.43	97.67 ± 0.43	97.72 ± 0.43	97.67 ± 0.43	97.65 ± 0.41	0.25 ± 0.01	97.66 ± 0.34	97.66 ± 0.34	97.66 ± 0.33	97.66 ± 0.34	97.66 ± 0.33	0.40 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	97.58 ± 0.43	97.58 ± 0.43	97.64 ± 0.35	97.58 ± 0.43	97.55 ± 0.47	0.27 ± 0.01	97.57 ± 0.31	97.57 ± 0.31	97.60 ± 0.32	97.57 ± 0.31	97.55 ± 0.31	0.42 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	97.59 ± 0.30	97.59 ± 0.30	97.62 ± 0.29	97.59 ± 0.30	97.58 ± 0.31	0.28 ± 0.01	97.52 ± 0.29	97.52 ± 0.29	97.54 ± 0.28	97.52 ± 0.29	97.50 ± 0.30	0.44 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	97.57 ± 0.37	97.57 ± 0.37	97.60 ± 0.38	97.57 ± 0.37	97.56 ± 0.36	0.30 ± 0.01	97.38 ± 0.47	97.38 ± 0.47	97.41 ± 0.43	97.38 ± 0.47	97.41 ± 0.46	0.47 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	97.63 ± 0.38	97.63 ± 0.38	97.65 ± 0.35	97.63 ± 0.38	97.63 ± 0.39	0.31 ± 0.01	97.56 ± 0.34	97.56 ± 0.34	97.57 ± 0.33	97.56 ± 0.34	97.57 ± 0.34	0.50 ± 0.01
3 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	97.36 ± 0.48	97.36 ± 0.48	97.41 ± 0.47	97.36 ± 0.48	97.35 ± 0.46	0.33 ± 0.01	97.30 ± 0.48	97.30 ± 0.48	97.41 ± 0.42	97.30 ± 0.48	97.24 ± 0.50	0.52 ± 0.01
3 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	97.62 ± 0.48	97.62 ± 0.48	97.65 ± 0.45	97.62 ± 0.48	97.61 ± 0.49	0.35 ± 0.01	97.52 ± 0.20	97.52 ± 0.20	97.54 ± 0.22	97.52 ± 0.20	97.50 ± 0.20	0.55 ± 0.01
3 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	97.46 ± 0.35	97.46 ± 0.35	97.47 ± 0.35	97.46 ± 0.35	97.46 ± 0.36	0.37 ± 0.01	97.50 ± 0.39	97.50 ± 0.39	97.51 ± 0.38	97.50 ± 0.39	97.51 ± 0.37	0.58 ± 0.01
128 Units/Layer			MisRoBÆRTa with LSTM Cells						MisRoBÆRTa with BiLSTM Cells					
BART Branch	RoBERTa Branch	Ensemble Branch	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)
1 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	97.17 ± 0.54	97.17 ± 0.54	97.22 ± 0.50	97.17 ± 0.54	97.17 ± 0.54	0.01 ± 0.00	97.23 ± 0.52	97.23 ± 0.52	97.26 ± 0.49	97.23 ± 0.52	97.25 ± 0.49	0.01 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	93.98 ± 0.97	93.98 ± 0.97	97.06 ± 0.84	93.98 ± 0.97	97.01 ± 0.92	0.02 ± 0.00	97.32 ± 0.27	97.32 ± 0.27	97.36 ± 0.24	97.32 ± 0.27	97.31 ± 0.29	0.03 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	97.04 ± 0.55	97.04 ± 0.55	97.10 ± 0.47	97.04 ± 0.55	97.04 ± 0.59	0.03 ± 0.00	97.19 ± 0.32	97.19 ± 0.32	97.21 ± 0.30	97.19 ± 0.32	97.20 ± 0.32	0.04 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	97.18 ± 0.61	97.18 ± 0.61	97.25 ± 0.60	97.18 ± 0.61	97.18 ± 0.58	0.03 ± 0.00	97.33 ± 0.42	97.33 ± 0.42	97.38 ± 0.44	97.33 ± 0.42	97.33 ± 0.42	0.06 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	97.13 ± 0.81	97.13 ± 0.81	97.19 ± 0.76	97.13 ± 0.81	97.14 ± 0.79	0.05 ± 0.00	97.37 ± 0.26	97.37 ± 0.26	97.37 ± 0.25	97.37 ± 0.26	97.38 ± 0.27	0.08 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	97.34 ± 0.48	97.34 ± 0.48	97.37 ± 0.47	97.34 ± 0.48	97.34 ± 0.47	0.06 ± 0.00	97.03 ± 0.86	97.03 ± 0.86	97.11 ± 0.72	97.03 ± 0.86	97.06 ± 0.80	0.10 ± 0.00
1 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	97.38 ± 0.65	97.38 ± 0.65	97.40 ± 0.64	97.38 ± 0.65	97.37 ± 0.64	0.07 ± 0.00	97.04 ± 0.53	97.04 ± 0.53	97.14 ± 0.44	97.04 ± 0.53	97.03 ± 0.54	0.12 ± 0.01
1 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	97.23 ± 0.41	97.23 ± 0.41	97.25 ± 0.38	97.23 ± 0.41	97.22 ± 0.43	0.08 ± 0.00	97.15 ± 0.42	97.15 ± 0.42	97.17 ± 0.40	97.15 ± 0.42	97.17 ± 0.42	0.14 ± 0.01
1 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	97.43 ± 0.40	97.43 ± 0.40	97.45 ± 0.41	97.43 ± 0.40	97.42 ± 0.38	0.09 ± 0.00	97.24 ± 0.28	97.24 ± 0.28	97.24 ± 0.28	97.24 ± 0.28	97.28 ± 0.28	0.16 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	97.55 ± 0.38	97.55 ± 0.38	97.63 ± 0.37	97.55 ± 0.38	97.51 ± 0.38	0.10 ± 0.00	97.58 ± 0.73	97.58 ± 0.73	97.55 ± 0.65	97.58 ± 0.73	97.58 ± 0.70	0.18 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	97.43 ± 0.35	97.53 ± 0.35	97.48 ± 0.36	97.43 ± 0.35	97.42 ± 0.34	0.11 ± 0.00	97.18 ± 0.46	97.18 ± 0.46	97.24 ± 0.40	97.18 ± 0.46	97.16 ± 0.48	0.19 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	97.35 ± 0.47	97.35 ± 0.47	97.40 ± 0.46	97.35 ± 0.47	97.34 ± 0.45	0.13 ± 0.00	97.36 ± 0.27	97.36 ± 0.27	97.37 ± 0.28	97.36 ± 0.27	97.35 ± 0.26	0.22 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	97.00 ± 0.97	97.00 ± 0.97	97.16 ± 0.79	97.00 ± 0.97	93.97 ± 0.97	0.14 ± 0.00	97.04 ± 0.54	97.04 ± 0.54	97.13 ± 0.44	97.04 ± 0.54	97.05 ± 0.51	0.23 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	97.36 ± 0.57	97.36 ± 0.57	97.39 ± 0.54	97.36 ± 0.57	97.36 ± 0.57	0.15 ± 0.00	97.25 ± 0.39	97.25 ± 0.39	97.29 ± 0.35	97.25 ± 0.39	97.26 ± 0.39	0.25 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	97.23 ± 0.47	97.23 ± 0.47	97.29 ± 0.43	97.23 ± 0.47	97.23 ± 0.50	0.16 ± 0.00	97.23 ± 0.21	97.23 ± 0.21	97.26 ± 0.21	97.23 ± 0.21	97.24 ± 0.18	0.28 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	97.31 ± 0.60	97.31 ± 0.60	97.37 ± 0.58	97.31 ± 0.60	97.30 ± 0.59	0.18 ± 0.01	97.24 ± 0.44	97.24 ± 0.44	97.32 ± 0.37	97.24 ± 0.44	97.21 ± 0.47	0.30 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	97.47 ± 0.51	97.47 ± 0.51	97.53 ± 0.49	97.47 ± 0.51	97.45 ± 0.50	0.19 ± 0.01	97.26 ± 0.17	97.26 ± 0.17	97.31 ± 0.17	97.26 ± 0.17	97.24 ± 0.16	0.32 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	97.26 ± 0.43	97.26 ± 0.43	97.28 ± 0.41	97.26 ± 0.43	97.29 ± 0.43	0.21 ± 0.01	97.23 ± 0.31	97.23 ± 0.31	97.28 ± 0.25	97.23 ± 0.31	97.23 ± 0.35	0.35 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	97.40 ± 0.58	97.40 ± 0.58	97.45 ± 0.55	97.40 ± 0.58	97.39 ± 0.57	0.22 ± 0.01	97.40 ± 0.35	97.40 ± 0.35	97.40 ± 0.34	97.40 ± 0.35	97.43 ± 0.34	0.37 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	97.37 ± 0.47	97.37 ± 0.47	97.42 ± 0.48	97.37 ± 0.47	97.38 ± 0.45	0.23 ± 0.01	97.33 ± 0.34	97.33 ± 0.34	97.38 ± 0.32	97.33 ± 0.34	97.32 ± 0.32	0.39 ± 0.01

Table 7. Cont.

128 Units/Layer			MisRoBÆRTa with LSTM Cells						MisRoBÆRTa with BiLSTM Cells					
BART Branch	RoBERTa Branch	Ensemble Branch	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)
3 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	97.41 ± 0.61	97.41 ± 0.61	97.44 ± 0.61	97.41 ± 0.61	97.41 ± 0.60	0.25 ± 0.01	97.35 ± 0.30	97.35 ± 0.30	97.37 ± 0.29	97.35 ± 0.30	97.35 ± 0.31	0.41 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	97.35 ± 0.41	97.35 ± 0.41	97.40 ± 0.40	97.35 ± 0.41	97.33 ± 0.40	0.26 ± 0.01	97.06 ± 0.79	97.06 ± 0.79	97.19 ± 0.57	97.06 ± 0.79	97.03 ± 0.85	0.43 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	97.25 ± 0.58	97.25 ± 0.58	97.29 ± 0.57	97.25 ± 0.58	97.27 ± 0.58	0.27 ± 0.01	97.35 ± 0.41	97.35 ± 0.41	97.37 ± 0.38	97.35 ± 0.41	97.33 ± 0.43	0.46 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	97.24 ± 0.60	97.24 ± 0.60	97.29 ± 0.59	97.24 ± 0.60	97.26 ± 0.56	0.29 ± 0.01	97.13 ± 0.32	97.13 ± 0.32	97.15 ± 0.30	97.13 ± 0.32	97.15 ± 0.31	0.48 ± 0.01
3 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	97.47 ± 0.59	97.47 ± 0.59	97.52 ± 0.57	97.47 ± 0.59	97.47 ± 0.58	0.30 ± 0.01	97.34 ± 0.29	97.34 ± 0.29	97.35 ± 0.29	97.34 ± 0.29	97.36 ± 0.28	0.51 ± 0.02
3 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	97.45 ± 0.53	97.45 ± 0.53	97.49 ± 0.52	97.45 ± 0.53	97.46 ± 0.51	0.32 ± 0.01	97.23 ± 0.44	97.23 ± 0.44	97.24 ± 0.43	97.23 ± 0.44	97.24 ± 0.41	0.53 ± 0.02
3 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	97.48 ± 0.47	97.48 ± 0.47	97.49 ± 0.47	97.48 ± 0.47	97.48 ± 0.48	0.34 ± 0.01	97.23 ± 0.50	97.23 ± 0.50	97.27 ± 0.45	97.23 ± 0.50	97.23 ± 0.48	0.57 ± 0.02

Table 8. MisRoBÆRTa ablation and sensitivity testing using the FAKENEWSCORPUS dataset presented in Section 5. (Note: bold text marks the overall best result).

32 Units/Layer			MisRoBÆRTa with LSTM Cells						MisRoBÆRTa with BiLSTM Cells					
BART Branch	RoBERTa Branch	Ensemble Branch	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)
1 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	91.38 ± 0.24	91.38 ± 0.24	91.55 ± 0.26	91.38 ± 0.24	91.38 ± 0.24	0.03 ± 0.00	91.91 ± 0.30	91.91 ± 0.30	92.16 ± 0.20	91.91 ± 0.30	91.91 ± 0.30	0.03 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	91.35 ± 0.36	91.35 ± 0.36	91.57 ± 0.32	91.35 ± 0.36	91.35 ± 0.36	0.06 ± 0.00	91.90 ± 0.21	91.90 ± 0.21	92.08 ± 0.20	91.90 ± 0.21	91.90 ± 0.21	0.06 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	90.87 ± 0.34	90.87 ± 0.34	91.16 ± 0.35	90.87 ± 0.34	90.87 ± 0.34	0.11 ± 0.00	91.61 ± 0.21	91.61 ± 0.21	91.79 ± 0.17	91.61 ± 0.21	91.61 ± 0.21	0.10 ± 0.01
1 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	90.90 ± 0.35	90.90 ± 0.35	91.13 ± 0.30	90.90 ± 0.35	90.90 ± 0.35	0.14 ± 0.01	91.62 ± 0.29	91.62 ± 0.29	91.84 ± 0.28	91.62 ± 0.29	91.62 ± 0.29	0.13 ± 0.01
1 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	90.87 ± 0.25	90.87 ± 0.25	91.13 ± 0.25	90.87 ± 0.25	90.87 ± 0.25	0.18 ± 0.01	91.60 ± 0.16	91.60 ± 0.16	91.75 ± 0.17	91.60 ± 0.16	91.60 ± 0.16	0.17 ± 0.01
1 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	90.60 ± 0.35	90.60 ± 0.35	90.79 ± 0.30	90.60 ± 0.35	90.60 ± 0.35	0.23 ± 0.01	91.45 ± 0.21	91.45 ± 0.21	91.60 ± 0.24	91.45 ± 0.21	91.45 ± 0.21	0.21 ± 0.01
1 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	90.64 ± 0.43	90.64 ± 0.43	90.77 ± 0.43	90.64 ± 0.43	90.64 ± 0.43	0.26 ± 0.01	91.51 ± 0.22	91.51 ± 0.22	91.70 ± 0.24	91.51 ± 0.22	91.51 ± 0.22	0.25 ± 0.01
1 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	90.83 ± 0.21	90.83 ± 0.21	91.09 ± 0.18	90.83 ± 0.21	90.83 ± 0.21	0.31 ± 0.01	91.60 ± 0.21	91.60 ± 0.21	91.78 ± 0.24	91.60 ± 0.21	91.59 ± 0.21	0.29 ± 0.01
1 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	90.56 ± 0.34	90.56 ± 0.34	90.79 ± 0.32	90.56 ± 0.34	90.56 ± 0.34	0.36 ± 0.02	91.32 ± 0.18	91.32 ± 0.18	91.53 ± 0.16	91.32 ± 0.18	91.32 ± 0.18	0.35 ± 0.02
2 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	91.98 ± 0.19	91.98 ± 0.19	92.17 ± 0.15	91.98 ± 0.19	91.98 ± 0.19	0.39 ± 0.02	92.08 ± 0.18	92.08 ± 0.18	92.25 ± 0.19	92.08 ± 0.18	92.08 ± 0.18	0.38 ± 0.02
2 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	91.60 ± 0.19	91.60 ± 0.19	91.78 ± 0.25	91.60 ± 0.19	91.60 ± 0.19	0.43 ± 0.02	91.80 ± 0.17	91.80 ± 0.17	91.94 ± 0.18	91.80 ± 0.17	91.80 ± 0.17	0.42 ± 0.02
2 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	91.19 ± 0.28	91.19 ± 0.28	91.37 ± 0.27	91.19 ± 0.28	91.19 ± 0.28	0.48 ± 0.02	91.53 ± 0.30	91.53 ± 0.30	91.67 ± 0.30	91.53 ± 0.30	91.53 ± 0.30	0.47 ± 0.02
2 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	91.73 ± 0.12	91.73 ± 0.12	91.90 ± 0.09	91.73 ± 0.12	91.73 ± 0.12	0.51 ± 0.02	91.88 ± 0.17	91.88 ± 0.17	92.07 ± 0.20	91.88 ± 0.17	91.88 ± 0.17	0.51 ± 0.02
2 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	91.35 ± 0.24	91.35 ± 0.24	91.51 ± 0.22	91.35 ± 0.24	91.35 ± 0.24	0.56 ± 0.02	91.71 ± 0.20	91.71 ± 0.20	91.86 ± 0.19	91.71 ± 0.20	91.71 ± 0.20	0.56 ± 0.02
2 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	90.92 ± 0.37	90.92 ± 0.37	91.15 ± 0.34	90.92 ± 0.37	90.92 ± 0.37	0.61 ± 0.02	91.20 ± 0.25	91.20 ± 0.25	91.32 ± 0.24	91.20 ± 0.25	91.20 ± 0.25	0.62 ± 0.03
2 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	91.55 ± 0.18	91.55 ± 0.18	91.68 ± 0.21	91.55 ± 0.18	91.55 ± 0.18	0.64 ± 0.02	91.56 ± 0.21	91.56 ± 0.21	91.75 ± 0.19	91.56 ± 0.21	91.56 ± 0.21	0.67 ± 0.03
2 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	91.10 ± 0.46	91.10 ± 0.46	91.23 ± 0.44	91.10 ± 0.46	91.10 ± 0.46	0.69 ± 0.03	91.49 ± 0.21	91.49 ± 0.21	91.58 ± 0.24	91.49 ± 0.21	91.49 ± 0.21	0.72 ± 0.03
2 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	90.69 ± 0.37	90.69 ± 0.37	90.92 ± 0.34	90.69 ± 0.37	90.69 ± 0.37	0.75 ± 0.03	91.19 ± 0.25	91.19 ± 0.25	91.35 ± 0.32	91.19 ± 0.25	91.19 ± 0.25	0.79 ± 0.03
3 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	91.74 ± 0.23	91.74 ± 0.23	91.92 ± 0.22	91.74 ± 0.23	91.74 ± 0.23	0.80 ± 0.03	91.92 ± 0.22	91.92 ± 0.22	92.06 ± 0.23	91.92 ± 0.22	91.92 ± 0.22	0.84 ± 0.04
3 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	91.39 ± 0.20	91.39 ± 0.20	91.58 ± 0.17	91.39 ± 0.20	91.39 ± 0.20	0.85 ± 0.03	91.72 ± 0.32	91.72 ± 0.32	91.91 ± 0.31	91.72 ± 0.32	91.72 ± 0.32	0.92 ± 0.04
3 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	90.86 ± 0.20	90.86 ± 0.20	91.06 ± 0.16	90.86 ± 0.20	90.86 ± 0.20	0.90 ± 0.04	91.48 ± 0.26	91.48 ± 0.26	91.64 ± 0.23	91.48 ± 0.26	91.48 ± 0.26	0.97 ± 0.04
3 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	91.62 ± 0.19	91.62 ± 0.19	91.82 ± 0.17	91.62 ± 0.19	91.62 ± 0.19	0.95 ± 0.04	91.84 ± 0.25	91.84 ± 0.25	92.00 ± 0.28	91.84 ± 0.25	91.84 ± 0.25	1.02 ± 0.04
3 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	91.16 ± 0.26	91.16 ± 0.26	91.31 ± 0.26	91.16 ± 0.26	91.16 ± 0.26	1.01 ± 0.04	91.51 ± 0.19	91.51 ± 0.19	91.66 ± 0.17	91.51 ± 0.19	91.51 ± 0.19	1.09 ± 0.05
3 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	90.36 ± 0.00	90.36 ± 0.00	90.79 ± 0.00	90.36 ± 0.00	90.36 ± 0.00	1.03 ± 0.04	91.25 ± 0.27	91.25 ± 0.27	91.49 ± 0.27	91.25 ± 0.27	91.25 ± 0.27	1.17 ± 0.06
3 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	91.23 ± 0.28	91.23 ± 0.28	91.37 ± 0.26	91.23 ± 0.28	91.23 ± 0.28	1.09 ± 0.04	91.50 ± 0.23	91.50 ± 0.23	91.63 ± 0.21	91.50 ± 0.23	91.50 ± 0.23	1.23 ± 0.06
3 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	90.77 ± 0.28	90.77 ± 0.28	90.92 ± 0.30	90.77 ± 0.28	90.77 ± 0.28	1.15 ± 0.04	91.32 ± 0.19	91.32 ± 0.19	91.47 ± 0.18	91.32 ± 0.19	91.32 ± 0.19	1.30 ± 0.06
3 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	88.50 ± 0.36	88.50 ± 0.36	88.87 ± 0.24	88.50 ± 0.36	88.50 ± 0.36	1.18 ± 0.05	91.22 ± 0.25	91.22 ± 0.25	91.41 ± 0.20	91.22 ± 0.25	91.22 ± 0.25	1.36 ± 0.08

Table 8. Cont.

64 Units/Layer			MisRoBÆRTa with LSTM Cells						MisRoBÆRTa with BiLSTM Cells					
BART Branch	RoBERTa Branch	Ensemble Branch	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)
1 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	92.05 ± 0.18	92.05 ± 0.18	92.23 ± 0.19	92.05 ± 0.18	92.05 ± 0.18	0.02 ± 0.00	92.26 ± 0.23	92.26 ± 0.23	92.40 ± 0.21	92.26 ± 0.23	92.26 ± 0.23	0.02 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	91.77 ± 0.29	91.77 ± 0.29	91.96 ± 0.25	91.77 ± 0.29	91.77 ± 0.29	0.04 ± 0.00	92.12 ± 0.19	92.12 ± 0.19	92.22 ± 0.19	92.12 ± 0.19	92.12 ± 0.19	0.05 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	91.72 ± 0.21	91.72 ± 0.21	91.90 ± 0.21	91.72 ± 0.21	91.72 ± 0.21	0.07 ± 0.00	91.90 ± 0.17	91.90 ± 0.17	92.07 ± 0.18	91.90 ± 0.17	91.90 ± 0.17	0.08 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	91.72 ± 0.28	91.72 ± 0.28	91.92 ± 0.29	91.72 ± 0.28	91.72 ± 0.28	0.09 ± 0.00	92.02 ± 0.20	92.02 ± 0.20	92.16 ± 0.18	92.02 ± 0.20	92.02 ± 0.20	0.11 ± 0.01
1 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	91.63 ± 0.13	91.63 ± 0.13	91.83 ± 0.13	91.63 ± 0.13	91.63 ± 0.13	0.11 ± 0.00	91.92 ± 0.28	91.92 ± 0.28	92.14 ± 0.21	91.92 ± 0.28	91.92 ± 0.28	0.14 ± 0.01
1 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	91.48 ± 0.35	91.48 ± 0.35	91.73 ± 0.28	91.48 ± 0.35	91.48 ± 0.35	0.15 ± 0.01	91.97 ± 0.15	91.97 ± 0.15	92.06 ± 0.15	91.97 ± 0.15	91.97 ± 0.15	0.18 ± 0.01
1 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	91.57 ± 0.23	91.57 ± 0.23	91.78 ± 0.20	91.57 ± 0.23	91.57 ± 0.23	0.17 ± 0.01	92.01 ± 0.16	92.01 ± 0.16	92.16 ± 0.19	92.01 ± 0.16	92.01 ± 0.16	0.21 ± 0.01
1 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	91.66 ± 0.16	91.66 ± 0.16	91.79 ± 0.18	91.66 ± 0.16	91.66 ± 0.16	0.19 ± 0.01	91.90 ± 0.21	91.90 ± 0.21	92.08 ± 0.22	91.90 ± 0.21	91.90 ± 0.21	0.25 ± 0.01
1 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	91.37 ± 0.26	91.37 ± 0.26	91.55 ± 0.27	91.37 ± 0.26	91.37 ± 0.26	0.23 ± 0.01	91.64 ± 0.22	91.64 ± 0.22	91.77 ± 0.24	91.64 ± 0.22	91.64 ± 0.22	0.29 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	92.30 ± 0.30	92.30 ± 0.30	92.47 ± 0.24	92.30 ± 0.30	92.30 ± 0.30	0.25 ± 0.01	92.41 ± 0.18	92.41 ± 0.18	92.57 ± 0.13	92.41 ± 0.18	92.41 ± 0.18	0.32 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	92.00 ± 0.35	92.00 ± 0.35	92.14 ± 0.32	92.00 ± 0.35	92.00 ± 0.35	0.28 ± 0.01	92.22 ± 0.17	92.22 ± 0.17	92.40 ± 0.16	92.22 ± 0.17	92.22 ± 0.17	0.35 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	91.94 ± 0.29	91.94 ± 0.29	92.06 ± 0.31	91.94 ± 0.29	91.94 ± 0.29	0.31 ± 0.01	92.12 ± 0.27	92.12 ± 0.27	92.24 ± 0.25	92.12 ± 0.27	92.12 ± 0.27	0.40 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	92.06 ± 0.21	92.06 ± 0.21	92.18 ± 0.23	92.06 ± 0.21	92.06 ± 0.21	0.33 ± 0.01	92.32 ± 0.23	92.32 ± 0.23	92.49 ± 0.24	92.32 ± 0.23	92.32 ± 0.23	0.43 ± 0.02
2 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	91.83 ± 0.25	91.83 ± 0.25	91.99 ± 0.25	91.83 ± 0.25	91.83 ± 0.25	0.36 ± 0.01	92.16 ± 0.25	92.16 ± 0.25	92.31 ± 0.15	92.16 ± 0.25	92.16 ± 0.25	0.47 ± 0.02
2 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	91.61 ± 0.19	91.61 ± 0.19	91.77 ± 0.21	91.61 ± 0.19	91.61 ± 0.19	0.40 ± 0.01	91.74 ± 0.25	91.74 ± 0.25	91.91 ± 0.26	91.74 ± 0.25	91.74 ± 0.25	0.52 ± 0.02
2 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	91.97 ± 0.27	91.97 ± 0.27	92.10 ± 0.24	91.97 ± 0.27	91.97 ± 0.27	0.43 ± 0.01	92.01 ± 0.25	92.01 ± 0.25	92.19 ± 0.16	92.01 ± 0.25	92.01 ± 0.25	0.56 ± 0.02
2 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	91.81 ± 0.31	91.81 ± 0.31	91.97 ± 0.31	91.81 ± 0.31	91.81 ± 0.31	0.46 ± 0.02	91.81 ± 0.17	91.81 ± 0.17	91.96 ± 0.23	91.81 ± 0.17	91.81 ± 0.17	0.60 ± 0.02
2 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	91.47 ± 0.26	91.47 ± 0.26	91.64 ± 0.24	91.47 ± 0.26	91.47 ± 0.26	0.50 ± 0.02	91.51 ± 0.43	91.51 ± 0.43	91.70 ± 0.41	91.51 ± 0.43	91.51 ± 0.43	0.66 ± 0.02
3 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	92.29 ± 0.26	92.29 ± 0.26	92.43 ± 0.21	92.29 ± 0.26	92.29 ± 0.26	0.52 ± 0.02	92.19 ± 0.31	92.19 ± 0.31	92.31 ± 0.34	92.19 ± 0.31	92.19 ± 0.31	0.70 ± 0.02
3 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	91.90 ± 0.21	91.90 ± 0.21	92.09 ± 0.16	91.90 ± 0.21	91.90 ± 0.21	0.56 ± 0.02	91.97 ± 0.21	91.97 ± 0.21	92.09 ± 0.20	91.97 ± 0.21	91.97 ± 0.21	0.74 ± 0.03
3 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	91.85 ± 0.28	91.85 ± 0.28	91.98 ± 0.26	91.85 ± 0.28	91.85 ± 0.28	0.60 ± 0.02	91.91 ± 0.15	91.91 ± 0.15	92.06 ± 0.17	91.91 ± 0.15	91.91 ± 0.15	0.80 ± 0.03
3 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	92.08 ± 0.19	92.08 ± 0.19	92.20 ± 0.19	92.08 ± 0.19	92.08 ± 0.19	0.63 ± 0.02	92.16 ± 0.23	92.16 ± 0.23	92.31 ± 0.19	92.16 ± 0.23	92.16 ± 0.23	0.84 ± 0.03
3 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	91.73 ± 0.27	91.73 ± 0.27	91.93 ± 0.27	91.73 ± 0.27	91.73 ± 0.27	0.67 ± 0.02	91.91 ± 0.13	91.91 ± 0.13	92.03 ± 0.16	91.90 ± 0.13	91.91 ± 0.13	0.92 ± 0.03
3 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	91.43 ± 0.16	91.43 ± 0.16	91.63 ± 0.16	91.43 ± 0.16	91.43 ± 0.16	0.70 ± 0.02	91.59 ± 0.07	91.59 ± 0.07	91.71 ± 0.11	91.59 ± 0.07	91.59 ± 0.07	0.93 ± 0.04
3 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	91.92 ± 0.22	91.92 ± 0.22	92.03 ± 0.22	91.92 ± 0.22	91.92 ± 0.22	0.73 ± 0.02	91.97 ± 0.23	91.97 ± 0.23	92.13 ± 0.24	91.97 ± 0.23	91.97 ± 0.23	0.98 ± 0.04
3 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	91.61 ± 0.28	91.61 ± 0.28	91.77 ± 0.27	91.61 ± 0.28	91.61 ± 0.28	0.77 ± 0.03	91.80 ± 0.16	91.80 ± 0.16	91.93 ± 0.18	91.81 ± 0.16	91.80 ± 0.16	1.04 ± 0.04
3 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	89.81 ± 0.21	89.81 ± 0.21	89.91 ± 0.27	89.81 ± 0.21	89.81 ± 0.21	0.79 ± 0.03	85.53 ± 0.11	85.53 ± 0.11	85.64 ± 0.17	85.53 ± 0.11	85.53 ± 0.11	1.08 ± 0.04
128 Units/Layer			MisRoBÆRTa with LSTM Cells						MisRoBÆRTa with BiLSTM Cells					
BART Branch	RoBERTa Branch	Ensemble Branch	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)
1 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	92.16 ± 0.27	92.16 ± 0.27	92.35 ± 0.22	92.16 ± 0.27	92.16 ± 0.27	0.02 ± 0.00	92.20 ± 0.30	92.20 ± 0.30	92.35 ± 0.29	92.20 ± 0.30	92.20 ± 0.30	0.03 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	92.11 ± 0.37	92.11 ± 0.37	92.26 ± 0.35	92.11 ± 0.37	92.11 ± 0.37	0.03 ± 0.00	92.23 ± 0.14	92.23 ± 0.14	92.37 ± 0.13	92.23 ± 0.14	92.23 ± 0.14	0.05 ± 0.00
1 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	91.95 ± 0.29	91.95 ± 0.29	92.14 ± 0.23	91.95 ± 0.29	91.95 ± 0.29	0.05 ± 0.00	92.04 ± 0.25	92.04 ± 0.25	92.17 ± 0.24	92.04 ± 0.25	92.04 ± 0.25	0.09 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	92.02 ± 0.29	92.02 ± 0.29	92.17 ± 0.24	92.02 ± 0.29	92.02 ± 0.29	0.07 ± 0.00	92.17 ± 0.26	92.17 ± 0.26	92.36 ± 0.23	92.17 ± 0.26	92.17 ± 0.26	0.11 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	91.87 ± 0.24	91.87 ± 0.24	91.98 ± 0.26	91.87 ± 0.24	91.87 ± 0.24	0.09 ± 0.00	91.94 ± 0.23	91.94 ± 0.23	92.03 ± 0.25	91.94 ± 0.23	91.94 ± 0.23	0.15 ± 0.00
1 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	91.56 ± 0.32	91.56 ± 0.32	91.72 ± 0.32	91.56 ± 0.32	91.56 ± 0.32	0.11 ± 0.00	91.62 ± 0.27	91.62 ± 0.27	91.76 ± 0.27	91.62 ± 0.27	91.62 ± 0.27	0.18 ± 0.00
1 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	91.86 ± 0.30	91.86 ± 0.30	92.01 ± 0.31	91.86 ± 0.30	91.86 ± 0.30	0.13 ± 0.00	91.97 ± 0.18	91.97 ± 0.18	92.09 ± 0.18	91.97 ± 0.18	91.97 ± 0.18	0.21 ± 0.00
1 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	91.58 ± 0.52	91.58 ± 0.52	91.74 ± 0.51	91.58 ± 0.52	91.58 ± 0.52	0.15 ± 0.00	91.53 ± 0.28	91.53 ± 0.28	91.65 ± 0.33	91.53 ± 0.28	91.53 ± 0.28	0.25 ± 0.00

Table 8. Cont.

128 Units/Layer			MisRoBÆRTa with LSTM Cells						MisRoBÆRTa with BiLSTM Cells					
BART Branch	RoBERTa Branch	Ensemble Branch	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)	Accuracy	Precision Micro	Precision Macro	Recall Micro	Recall Macro	Execution Time (Hours)
1 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	91.09 ± 0.43	91.09 ± 0.43	91.21 ± 0.42	91.09 ± 0.43	91.09 ± 0.43	0.17 ± 0.00	91.07 ± 0.32	91.07 ± 0.32	91.17 ± 0.34	91.07 ± 0.32	91.07 ± 0.32	0.29 ± 0.00
2 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	92.50 ± 0.38	92.50 ± 0.38	92.67 ± 0.35	92.50 ± 0.38	92.50 ± 0.38	0.19 ± 0.00	92.50 ± 0.26	92.50 ± 0.26	92.69 ± 0.21	92.50 ± 0.26	92.50 ± 0.26	0.32 ± 0.00
2 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	92.30 ± 0.28	92.30 ± 0.28	92.46 ± 0.25	92.30 ± 0.28	92.30 ± 0.28	0.21 ± 0.01	92.35 ± 0.26	92.35 ± 0.26	92.52 ± 0.22	92.35 ± 0.26	92.35 ± 0.26	0.35 ± 0.01
2 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	92.04 ± 0.37	92.04 ± 0.37	92.20 ± 0.38	92.04 ± 0.37	92.04 ± 0.37	0.24 ± 0.00	92.11 ± 0.25	92.11 ± 0.25	92.20 ± 0.21	92.11 ± 0.25	92.11 ± 0.25	0.39 ± 0.00
2 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	92.36 ± 0.31	92.36 ± 0.31	92.48 ± 0.25	92.36 ± 0.31	92.36 ± 0.31	0.26 ± 0.00	92.40 ± 0.20	92.40 ± 0.20	92.52 ± 0.21	92.40 ± 0.20	92.40 ± 0.20	0.42 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	92.05 ± 0.23	92.05 ± 0.23	92.17 ± 0.21	92.05 ± 0.23	92.05 ± 0.23	0.28 ± 0.00	92.19 ± 0.25	92.19 ± 0.25	92.33 ± 0.22	92.19 ± 0.25	92.19 ± 0.25	0.46 ± 0.01
2 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	91.70 ± 0.47	91.70 ± 0.47	91.87 ± 0.45	91.70 ± 0.47	91.70 ± 0.47	0.31 ± 0.01	91.69 ± 0.35	91.69 ± 0.35	91.88 ± 0.40	91.69 ± 0.35	91.69 ± 0.35	0.51 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	92.13 ± 0.28	92.13 ± 0.28	92.36 ± 0.27	92.13 ± 0.28	92.13 ± 0.28	0.33 ± 0.01	92.16 ± 0.28	92.16 ± 0.28	92.33 ± 0.27	92.16 ± 0.28	92.16 ± 0.28	0.54 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	91.94 ± 0.32	91.94 ± 0.32	92.06 ± 0.33	91.94 ± 0.32	91.94 ± 0.32	0.35 ± 0.01	91.69 ± 0.32	91.69 ± 0.32	91.79 ± 0.37	91.69 ± 0.32	91.69 ± 0.32	0.58 ± 0.01
2 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	91.18 ± 0.25	91.18 ± 0.25	91.30 ± 0.27	91.18 ± 0.25	91.18 ± 0.25	0.38 ± 0.01	91.18 ± 0.23	91.18 ± 0.23	91.29 ± 0.24	91.18 ± 0.23	91.18 ± 0.23	0.63 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	1 × [Bi]LSTM	92.38 ± 0.25	92.38 ± 0.25	92.50 ± 0.22	92.38 ± 0.25	92.38 ± 0.25	0.40 ± 0.01	92.34 ± 0.21	92.34 ± 0.21	92.45 ± 0.23	92.34 ± 0.21	92.34 ± 0.21	0.67 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	2 × [Bi]LSTM	92.13 ± 0.30	92.13 ± 0.30	92.23 ± 0.30	92.13 ± 0.30	92.13 ± 0.30	0.43 ± 0.01	92.16 ± 0.27	92.16 ± 0.27	92.26 ± 0.26	92.16 ± 0.27	92.16 ± 0.27	0.71 ± 0.01
3 × [Bi]LSTM	1 × [Bi]LSTM	3 × [Bi]LSTM	91.87 ± 0.26	91.87 ± 0.26	92.13 ± 0.24	91.87 ± 0.26	91.87 ± 0.26	0.45 ± 0.01	92.06 ± 0.23	92.06 ± 0.23	92.16 ± 0.26	92.06 ± 0.23	92.06 ± 0.23	0.76 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	1 × [Bi]LSTM	92.28 ± 0.28	92.28 ± 0.28	92.42 ± 0.26	92.28 ± 0.28	92.28 ± 0.28	0.48 ± 0.01	92.22 ± 0.17	92.22 ± 0.17	92.37 ± 0.11	92.22 ± 0.17	92.22 ± 0.17	0.80 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	2 × [Bi]LSTM	92.04 ± 0.25	92.04 ± 0.25	92.20 ± 0.31	92.04 ± 0.25	92.04 ± 0.25	0.50 ± 0.01	92.04 ± 0.25	92.04 ± 0.25	92.17 ± 0.26	92.04 ± 0.25	92.04 ± 0.25	0.84 ± 0.01
3 × [Bi]LSTM	2 × [Bi]LSTM	3 × [Bi]LSTM	91.14 ± 0.21	91.14 ± 0.21	91.19 ± 0.24	91.14 ± 0.21	91.14 ± 0.21	0.52 ± 0.01	91.30 ± 0.11	91.30 ± 0.11	91.43 ± 0.12	91.30 ± 0.11	91.30 ± 0.11	0.88 ± 0.01
3 × [Bi]LSTM	3 × [Bi]LSTM	1 × [Bi]LSTM	92.19 ± 0.23	92.19 ± 0.23	92.31 ± 0.25	92.19 ± 0.23	92.19 ± 0.23	0.55 ± 0.01	92.19 ± 0.18	92.19 ± 0.18	92.27 ± 0.21	92.19 ± 0.18	92.19 ± 0.18	0.92 ± 0.02
3 × [Bi]LSTM	3 × [Bi]LSTM	2 × [Bi]LSTM	92.10 ± 0.13	92.10 ± 0.13	92.19 ± 0.12	92.10 ± 0.13	92.10 ± 0.13	0.58 ± 0.02	91.57 ± 0.32	91.57 ± 0.32	91.70 ± 0.35	91.57 ± 0.32	91.57 ± 0.32	0.97 ± 0.02
3 × [Bi]LSTM	3 × [Bi]LSTM	3 × [Bi]LSTM	91.01 ± 0.09	91.01 ± 0.09	91.09 ± 0.08	91.01 ± 0.09	91.01 ± 0.09	0.60 ± 0.02	91.03 ± 0.17	91.03 ± 0.17	91.14 ± 0.20	91.03 ± 0.17	91.03 ± 0.17	1.02 ± 0.02

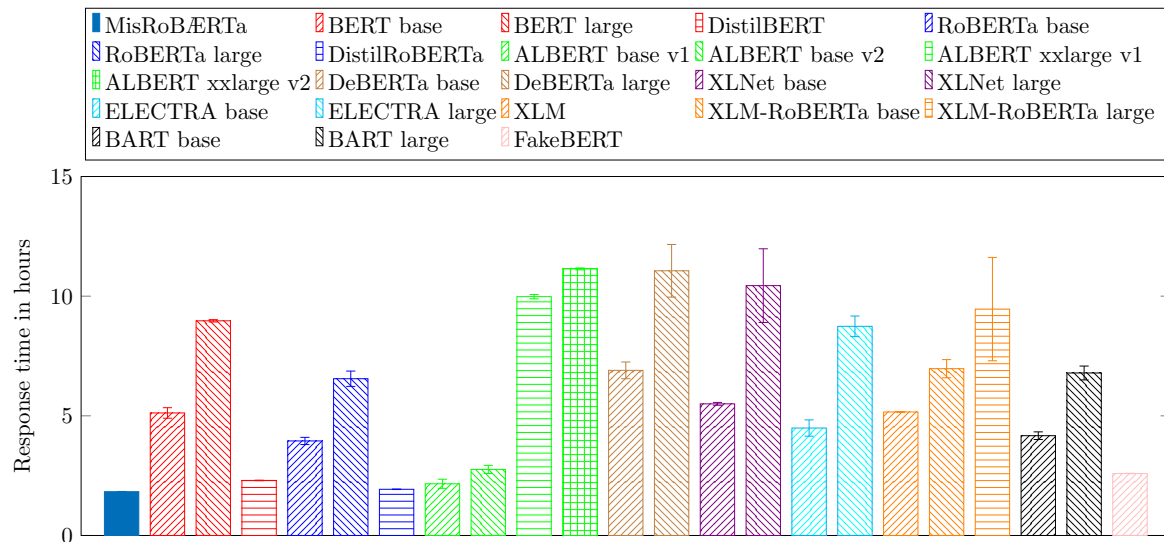


Figure 2. Performance time for the tested architectures.

6. Discussion

The misinformation detection problem is a complex task that requires large datasets to accurately determine news content veracity. We found that in the current literature, many authors utilize relatively small datasets for this task, e.g., datasets with ~ 1000 articles in [19], $\sim 10,000$ in [25], which is negatively reflected in the accuracy obtained by the tested models, especially when it comes to BERT-based models. Surely, there are authors who test their solutions using larger corpora, e.g., $\sim 80,000$ articles in [28]. In some cases, the dataset consists of short news statements, which in turn affect the accuracy considerably [18]. Furthermore, the current research articles only address the misinformation detection problem using either a binary approach (the news article's veracity is either true or false [8,19,28]) or levels of veracity [23]. Thus, none of the current approaches uses a large dataset that contains fake news labeled with different categories. We address this shortcoming by using a large dataset containing 100,000 news articles labeled with 10 classes.

The size of the corpus, number of tokens, the used vocabulary, the frequent unigrams, and topics play an important role in the overall and per class detection task. Through fine-tuning, we include these dimensions in our models, thus improving the classification performance.

In the misinformation detection task, contextual information plays an important role that impacts the accuracy of determining hidden patterns within the text. Thus, using transformers and transfer learning, the models better discriminate between different types of misinformation as shown by the recall score of over 85%. By adding autoregression, sequence-to-sequence with a bidirectional encoder, and a left-right decoder, BART large achieves the overall best performance among the basic architecture with transformers.

For BERT, RoBERTa, DeBERTa, XLNet, ELECTRA, and XLM-RoBERTa, the base models manage to better generalize than the large ones, while for BART and ALBERT, the reverse happens. Thus, we can conclude that BERT, RoBERTa, DeBERTa, XLNet, ELECTRA, and XLM-RoBERTa manage to extract correctly the hidden context within the misinformation texts and better discriminate between the labels with a smaller network size and number of parameters. The autoregressive method to learn bidirectional contexts employed by XLNet and BART improves the misinformation detection performance as the network and number of parameters increases. We can conclude that, for XLNet base, although extracting less information from context than large, it achieves better results through the autoregressive method. Thus, both XLNet base and BART (base and large) extract more information from context using the autoregressive method.

XLM, a multilingual model, performs rather poorly. This is somehow expected as the best performing multilingual system scored below the worst performing native system [55].

Although distillation is a good technique to maintain classification accuracy by lowering resource consumption, DistilBERT and ALBERT models do not provide better results than the base models. As ALBERT also reduces BERT's number of parameters, it still needs a large network to achieve good results. Furthermore, ALBERT performance significantly improves when dropout is used, which also lowers execution performance. These models can still be used in low resource environments, offering an adequate level of accuracy. We can conclude that by reducing both the number of parameters and the network size, the classification performance is decreased for these models.

DistilRoBERTa achieves the best accuracy (91.32%) among the distilled versions and the overall smallest runtime (1.93 h for fine-tuning and training). Its accuracy falls short with only 0.63% than the model with the highest accuracy, BART large. In conclusion, DistilRoBERTa model can be a viable option when the number of resources is limited, and both fine-tuning and training should be conducted with the least amount of time.

MisRoBERTa is a BiLSTM-CNN deep learning architecture that utilizes as input both RoBERTa and BART and sentence embeddings. This novel architecture manages to better preserve context than the transformer models, as it employs context-aware and future selection layers besides context-aware embeddings. MisRoBERTa outperforms the BART large model w.r.t. classification performance and DistilRoBERTa model w.r.t. execution time.

Finally, we would like to conclude the discussion of our results by answering one final question: *is misinformation detection a solved problem?* In short, no. Although the presented transformer-based methods and the novel architecture MisRoBERTa show their effectiveness by obtaining high accuracy using linguistic features, there remain a lot of challenges that, if tackled, will improve performance and help users make informed decisions about the content of their news feeds. Some of the challenges we identified while surveying the related work for the task of misinformation detection that require a shift from a model-centric perspective to a data-centric perspective are:

- (1) Multi-modal (i.e., visual, audio, network propagation, and immunization, etc.) techniques are rarely used; this is due to a lack of complete, high quality, and well labeled corpora and the tedious work that is required to correctly annotate such datasets;
- (2) Source verification is rarely taken into consideration for misinformation detection [56], this is due to the instability of links on the internet;
- (3) Author credibility should be an important factor that is not really discussed or weighted in the feature selection for the models;
- (4) The perpetual need to retrain or fine-tune models to capture the subtle shifts that appear over time in news articles that spread misinformation.

7. Conclusions

In this paper, we present a performance evaluation of transformers on the task of misinformation detection. We employed a large real-world dataset for our experiments and addressed two shortcomings in the current research: (1) Increasing the size of the dataset from small to large; and (2) Moving the focus of fake news detection from binary to multi-class classification. We present and discuss our findings in detail. Furthermore, the results show that the accuracy of transformers on the misinformation detection task is significantly influenced by the method employed to learn the context, dataset size, and vocabulary.

We propose MisRoBERTa, a new transformer-based deep learning architecture for misinformation detection. The proposed architecture used, as input, BART and RoBERTa sentence embeddings and a BiLSTM-CNN architecture for each embedding. The output of each BiLSTM-CNN was concatenated and passed through another BiLSTM-CNN to obtain the results. We conclude that the proposed architecture outperforms the other employed classification models w.r.t. (1) classification performance, obtaining an average accuracy of 92.50%; and (2) runtime, with an average execution time for training of 1.83 h.

We respond to the research questions as follows. The overall best accuracy among the transformer models is obtained by BART (91.94%), while DistilRoBERTa obtains the

best accuracy (91.32%) in the least amount of time required for fine-tuning and training, i.e., 1.93 h on average. MisRoBERTa outperforms both of these models' w.r.t. accuracies (92.50%) and performance times (1.83 h). We empirically observed that between BART base and large, the difference in performance is insignificant, while the difference in runtime is of almost 2.62 h on average. Moreover, some improved models of BERT slightly outperform the base model, i.e., DeBERTa and RoBERTa, while most fall short. Furthermore, multilingual models, such as XLM, do not obtain better accuracy than BERT base.

In future research, our solution can be enhanced with sentiment analysis techniques that will determine the polarity and strength of the sentiments expressed in each news article. Another research direction is to add global context through topic modeling [57,58].

Author Contributions: Conceptualization, C.-O.T. and E.-S.A.; methodology, C.-O.T. and E.-S.A.; software, C.-O.T. and E.-S.A.; validation, C.-O.T. and E.-S.A.; formal analysis, C.-O.T. and E.-S.A.; investigation, C.-O.T. and E.-S.A.; resources, C.-O.T. and E.-S.A.; data curation, C.-O.T. and E.-S.A.; writing—original draft preparation, C.-O.T. and E.-S.A.; writing—review and editing, C.-O.T. and E.-S.A.; visualization, C.-O.T. and E.-S.A.; supervision, C.-O.T. and E.-S.A.; project administration, C.-O.T. and E.-S.A.; funding acquisition, C.-O.T. and E.-S.A. All authors have read and agreed to the published version of the manuscript.

Funding: This research received no external funding.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable.

Data Availability Statement: The original datasets used in this study are publicly available as follows: FakeNewsCorpus <https://github.com/several27/FakeNewsCorpus> (accessed on 9 October 2021); Kaggle Fake News Dataset <https://www.kaggle.com/c/fake-news/data> (accessed on 9 October 2021).

Acknowledgments: The publication of this paper is supported by the University Politehnica of Bucharest through the PubArt program and the work has been partly funded by EU CEF grant number 2394203 (NORDIS - NORDic observatory for digital media and information DISorders).

Conflicts of Interest: The authors declare no conflict of interest.

Abbreviations

The following abbreviations are used in this manuscript:

LSTM	long short-term memory
BiLSTM	bidirectional long short-term memory
CNN	convolutional neural network
BERT	bidirectional encoder representations from transformers
DistilBERT	distilled BERT
RoBERTa	robustly optimized BERT pre-training approach
DistilRoBERTa	distilled RoBERTa
ALBERT	a lite BERT
DeBERTa	decoding-enhanced BERT with disentangled attention
ELECTRA	efficiently learning an encoder that classifies token replacements Accurately
XLM	cross-lingual language model
BART	bidirectional and autoregressive transformer

References

1. Ruths, D. The misinformation machine. *Science* **2019**, *363*, 348. [CrossRef] [PubMed]
2. Shu, K.; Wang, S.; Le, T.; Lee, D.; Liu, H. Deep Headline Generation for Clickbait Detection. In Proceedings of the IEEE International Conference on Data Mining, Singapore, 17–20 November 2018. [CrossRef]
3. Ilie, V.I.; Truica, C.O.; Apostol, E.S.; Paschke, A. Context-Aware Misinformation Detection: A Benchmark of Deep Learning Architectures Using Word Embeddings. *IEEE Access* **2021**, *9*, 162122–162146. [CrossRef]
4. Zannettou, S.; Sirivianos, M.; Blackburn, J.; Kourtellis, N. The web of false information: Rumors, fake news, hoaxes, clickbait, and various other shenanigans. *J. Data Inf. Qual.* **2019**, *11*, 1–37. [CrossRef]

5. Bovet, A.; Makse, H.A. Influence of fake news in Twitter during the 2016 US presidential election. *Nat. Commun.* **2019**, *10*, 7. [CrossRef]
6. Marco-Franco, J.E.; Pita-Barros, P.; Vivas-Orts, D.; González-de-Julián, S.; Vivas-Consuelo, D. COVID-19, Fake News, and Vaccines: Should Regulation Be Implemented? *Int. J. Environ. Res. Public Health* **2021**, *18*, 744. [CrossRef]
7. Commission, E. Fighting Disinformation. 2021. Available online: https://ec.europa.eu/info/live-work-travel-eu/health/coronavirus-response/fighting-disinformation_en (accessed on 6 September 2021).
8. Kaliyar, R.K.; Goswami, A.; Narang, P. FakeBERT: Fake news detection in social media with a BERT-based deep learning approach. *Multimed. Tools Appl.* **2021**, *80*, 11765–11788. [CrossRef]
9. Conroy, N.K.; Rubin, V.L.; Chen, Y. Automatic deception detection: Methods for finding fake news. *Proc. Assoc. Inf. Sci. Technol.* **2015**, *52*, 1–4. [CrossRef]
10. Gravanis, G.; Vakali, A.; Diamantaras, K.; Karadais, P. Behind the cues: A benchmarking study for fake news detection. *Expert Syst. Appl.* **2019**, *128*, 201–213. [CrossRef]
11. Kaliyar, R.K.; Goswami, A.; Narang, P.; Sinha, S. FNDNet—A deep convolutional neural network for fake news detection. *Cogn. Syst. Res.* **2020**, *61*, 32–44. [CrossRef]
12. Pennington, J.; Socher, R.; Manning, C. GloVe: Global Vectors for Word Representation. In Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP), Doha, Qatar, 25–29 October 2014; pp. 1532–1543. [CrossRef]
13. Li, Q.; Hu, Q.; Lu, Y.; Yang, Y.; Cheng, J. Multi-level word features based on CNN for fake news detection in cultural communication. *Pers. Ubiquitous Comput.* **2020**, *24*, 259–272. [CrossRef]
14. Mikolov, T.; Chen, K.; Corrado, G.; Dean, J. Efficient estimation of word representations in vector space. In Proceedings of the International Conference on Learning Representations, Scottsdale, AZ, USA, 2–4 May 2013; pp. 1–12.
15. Mishra, R.; Setty, V. Sadhan: Hierarchical attention networks to learn latent aspect embeddings for fake news detection. In Proceedings of the 2019 ACM SIGIR International Conference on Theory of Information Retrieval, Santa Clara, CA, USA, 2–5 October 2019; pp. 197–204. [CrossRef]
16. Trueman, T.E.; Kumar, A.; Narayanasamy, P.; Vidya, J. Attention-based C-BiLSTM for fake news detection. *Appl. Soft Comput.* **2021**, *110*, 107600. [CrossRef]
17. Liu, Y.; Ott, M.; Goyal, N.; Du, J.; Joshi, M.; Chen, D.; Levy, O.; Lewis, M.; Zettlemoyer, L.; Stoyanov, V. RoBERTa: A Robustly Optimized BERT Pretraining Approach. *arXiv* **2019**, arXiv:1907.11692.
18. Liu, C.; Wu, X.; Yu, M.; Li, G.; Jiang, J.; Huang, W.; Lu, X. A Two-Stage Model Based on BERT for Short Fake News Detection. In Proceedings of the International Conference on Knowledge Science, Engineering and Management, Athens, Greece, 28–30 August 2019; pp. 172–183. [CrossRef]
19. Kurasinski, L.; Mihailescu, R.C. Towards Machine Learning Explainability in Text Classification for Fake News Detection. In Proceedings of the IEEE International Conference on Machine Learning and Applications, Miami, FL, USA, 14–17 December 2020; pp. 775–781. [CrossRef]
20. Jwa, H.; Oh, D.; Park, K.; Kang, J.; Lim, H. exBAKE: Automatic Fake News Detection Model Based on Bidirectional Encoder Representations from Transformers (BERT). *Appl. Sci.* **2019**, *9*, 4062. [CrossRef]
21. Mersinias, M.; Afantenos, S.; Chalkiadakis, G. CLFD: A Novel Vectorization Technique and Its Application in Fake News Detection. In Proceedings of the Language Resources and Evaluation Conference, Marseille, France, 11–16 May 2020; pp. 3475–3483.
22. Shaar, S.; Babulkov, N.; Da San Martino, G.; Nakov, P. That is a Known Lie: Detecting Previously Fact-Checked Claims. In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, Online, 5–10 July 2020; pp. 3607–3618. [CrossRef]
23. Kula, S.; Choraś, M.; Kozik, R. Application of the BERT-Based Architecture in Fake News Detection. In Proceedings of the Conference on Complex, Intelligent, and Software Intensive Systems, Asan, Korea, 1–3 July 2020; pp. 239–249. [CrossRef]
24. Shu, K.; Zheng, G.; Li, Y.; Mukherjee, S.; Awadallah, A.H.; Ruston, S.; Liu, H. Leveraging multi-source weak social supervision for early detection of fake news. *arXiv* **2020**, arXiv:2004.01732.
25. Gautam, A.; Venkatesh, V.; Masud, S. Fake News Detection System using XLNet model with Topic Distributions: CON-STRANT@AAAI2021 Shared Task. In Proceedings of the CONSTRAINT Shared Task in AAAI-2021, Virtual Event, 8 February 2021; pp. 1–12. [CrossRef]
26. Lan, Z.; Chen, M.; Goodman, S.; Gimpel, K.; Sharma, P.; Soricut, R. ALBERT: A Lite BERT for Self-supervised Learning of Language Representations. In Proceedings of the International Conference on Learning Representations, Addis Ababa, Ethiopia, 26–30 April 2020; pp. 1–17.
27. Tian, L.; Zhang, X.; Peng, M. FakeFinder: Twitter Fake News Detection on Mobile. In Proceedings of the Companion Proceedings of the Web Conference 2020, Taipei, Taiwan, 20–24 April 2020; pp. 79–80. [CrossRef]
28. Khan, J.Y.; Khondaker, M.T.I.; Afroz, S.; Uddin, G.; Iqbal, A. A benchmark study of machine learning models for online fake news detection. *Mach. Learn. Appl.* **2021**, *4*, 100032. [CrossRef]
29. Peters, M.E.; Neumann, M.; Iyyer, M.; Gardner, M.; Clark, C.; Lee, K.; Zettlemoyer, L. Deep Contextualized Word Representations. In Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers); Association for Computational Linguistics: New Orleans, LA, USA, 2018; pp. 2227–2237. [CrossRef]

30. Devlin, J.; Chang, M.W.; Lee, K.; Toutanova, K. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics, Minneapolis, MN, USA, 2–7 June 2019; pp. 4171–4186. [\[CrossRef\]](#)
31. Sanh, V.; Debut, L.; Chaumond, J.; Wolf, T. DistilBERT, a distilled version of BERT: Smaller, faster, cheaper and lighter. In Proceedings of the Workshop on Energy Efficient Machine Learning and Cognitive Computing, Vancouver, BC, Canada, 13 December 2019; pp. 1–5.
32. Sajjad, H.; Dalvi, F.; Durrani, N.; Nakov, P. Poor Man’s BERT: Smaller and Faster Transformer Models. *arXiv* **2020**, arXiv:2004.03844.
33. He, P.; Liu, X.; Gao, J.; Chen, W. DeBERTa: Decoding-enhanced BERT with Disentangled Attention. *arXiv* **2021**, arXiv:2006.03654.
34. Yang, Z.; Dai, Z.; Yang, Y.; Carbonell, J.; Salakhutdinov, R.; Le, Q.V. XLNet: Generalized Autoregressive Pretraining for Language Understanding. In Proceedings of the Advances in Neural Information Processing Systems, Vancouver, BC, Canada, 8–14 December 2019; pp. 5753–5763.
35. Dai, Z.; Yang, Z.; Yang, Y.; Carbonell, J.; Le, Q.; Salakhutdinov, R. Transformer-XL: Attentive Language Models beyond a Fixed-Length Context. In Proceedings of the Annual Meeting of the Association for Computational Linguistics, Florence, Italy, 28 July–2 August 2019; pp. 2978–2988. [\[CrossRef\]](#)
36. Clark, K.; Luong, M.T.; Le, Q.V.; Manning, C.D. ELECTRA: Pre-training Text Encoders as Discriminators Rather Than Generators. In Proceedings of the International Conference on Learning Representations, Addis Ababa, Ethiopia, 26–30 April 2020; pp. 1–17.
37. Conneau, A.; Lample, G. Cross-lingual Language Model Pretraining. In Proceedings of the Advances in Neural Information Processing Systems, Vancouver, BC, Canada, 8–14 December 2019; pp. 1–11.
38. Conneau, A.; Khandelwal, K.; Goyal, N.; Chaudhary, V.; Wenzek, G.; Guzmán, F.; Grave, E.; Ott, M.; Zettlemoyer, L.; Stoyanov, V. Unsupervised Cross-lingual Representation Learning at Scale. In Proceedings of the Annual Meeting of the Association for Computational Linguistics, Online, 5–10 July 2020; pp. 8440–8451. [\[CrossRef\]](#)
39. Lewis, M.; Liu, Y.; Goyal, N.; Ghazvininejad, M.; Mohamed, A.; Levy, O.; Stoyanov, V.; Zettlemoyer, L. BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension. In Proceedings of the Annual Meeting of the Association for Computational Linguistics, Online, 5–10 July 2020; pp. 7871–7880. [\[CrossRef\]](#)
40. Seide, F.; Li, G.; Chen, X.; Yu, D. Feature engineering in Context-Dependent Deep Neural Networks for conversational speech transcription. In Proceedings of the IEEE Workshop on Automatic Speech Recognition & Understanding, Waikoloa, HI, USA, 11–15 December 2011. [\[CrossRef\]](#)
41. Hochreiter, S.; Schmidhuber, J. Long Short-Term Memory. *Neural Comput.* **1997**, *9*, 1735–1780. [\[CrossRef\]](#)
42. Guo, B.; Zhang, C.; Liu, J.; Ma, X. Improving text classification with weighted word embeddings via a multi-channel TextCNN model. *Neurocomputing* **2019**, *363*, 366–374. [\[CrossRef\]](#)
43. Szpakowski, M. Fake News Corpus. 2021. Available online: <https://github.com/several27/FakeNewsCorpus> (accessed on 6 August 2021).
44. Geva, M.; Goldberg, Y.; Berant, J. Are We Modeling the Task or the Annotator? An Investigation of Annotator Bias in Natural Language Understanding Datasets. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), Hong Kong, China, 3–7 November 2019; pp. 1161–1166. [\[CrossRef\]](#)
45. Kuwatly, H.A.; Wich, M.; Groh, G. Identifying and Measuring Annotator Bias Based on Annotators’ Demographic Characteristics. In Proceedings of the Fourth Workshop on Online Abuse and Harms. Association for Computational Linguistics, Online, 20 November 2020; pp. 184–190. [\[CrossRef\]](#)
46. Jie, Z.; Xie, P.; Lu, W.; Ding, R.; Li, L. Better Modeling of Incomplete Annotations for Named Entity Recognition. In Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Minneapolis, MN, USA, 2–7 June 2019; pp. 729–734. [\[CrossRef\]](#)
47. Bird, S.; Loper, E.; Klein, E. *Natural Language Processing with Python*; O’Reilly Media Inc.: Sebastopol, CA, USA, 2009.
48. Grootendorst, M. PolyFuzz. 2021. Available online: <https://maartengr.github.io/PolyFuzz/> (accessed on 14 October 2021).
49. Arora, S.; Ge, R.; Moitra, A. Learning Topic Models—Going beyond SVD. In Proceedings of the Annual Symposium on Foundations of Computer Science, New Brunswick, NJ, USA, 20–23 October 2012; pp. 1–10.
50. Pedregosa, F.; Varoquaux, G.; Gramfort, A.; Michel, V.; Thirion, B.; Grisel, O.; Blondel, M.; Prettenhofer, P.; Weiss, R.; Dubourg, V.; et al. Scikit-learn: Machine Learning in Python. *J. Mach. Learn. Res.* **2011**, *12*, 2825–2830.
51. Rajapakse, T. SimpleTransformers. 2021. Available online: <https://simpletransformers.ai> (accessed on 12 October 2021).
52. Wolf, T.; Debut, L.; Sanh, V.; Chaumond, J.; Delangue, C.; Moi, A.; Cistac, P.; Rault, T.; Louf, R.; Funtowicz, M.; et al. Transformers: State-of-the-Art Natural Language Processing. In Proceedings of the Conference on Empirical Methods in Natural Language Processing, Punta Cana, Dominican Republic, 8–12 November 2020; pp. 38–45.
53. Abadi, M.; Agarwal, A.; Barham, P.; Brevdo, E.; Chen, Z.; Citro, C.; Corrado, G.S.; Davis, A.; Dean, J.; Devin, M.; et al. TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems. 2015. Available online: <https://www.tensorflow.org/> (accessed on 8 October 2021).
54. Chollet, F.; Jin, H.; Nozawa, K. Keras. 2015. Available online: <https://keras.io> (accessed on 8 October 2021).
55. Zampieri, M.; Nakov, P.; Rosenthal, S.; Atanasova, P.; Karadzhov, G.; Mubarak, H.; Derczynski, L.; Pitenis, Z.; Çağrı, Ç. SemEval-2020 Task 12: Multilingual Offensive Language Identification in Social Media (OffensEval 2020). In Proceedings of the Workshop on Semantic Evaluation, Online, 12–13 December 2020; pp. 1425–1447.

56. Parikh, S.B.; Atrey, P.K. Media-rich fake news detection: A survey. In Proceedings of the 2018 IEEE Conference on Multimedia Information Processing and Retrieval (MIPR), Miami, FL, USA, 10–12 April 2018; pp. 436–441. [\[CrossRef\]](#)
57. Mitroi, M.; Truică, C.O.; Apostol, E.S.; Florea, A.M. Sentiment Analysis using Topic-Document Embeddings. In Proceedings of the IEEE 16th International Conference on Intelligent Computer Communication and Processing, Cluj-Napoca, Romania, 3–5 September 2020; pp. 75–82. [\[CrossRef\]](#)
58. Truică, C.O.; Apostol, E.S.; Șerban, M.L.; Paschke, A. Topic-Based Document-Level Sentiment Analysis Using Contextual Cues. *Mathematics* **2021**, *9*, 2722. [\[CrossRef\]](#)